

# COMP0188

# Deep Representation and Learning

Joshua Spear  
[joshua.spear.21@ucl.ac.uk](mailto:joshua.spear.21@ucl.ac.uk)

# Today

- Coursework 1
- Invariance and equivariance
- CNNs
- Inducing invariance and equivariance through data augmentations

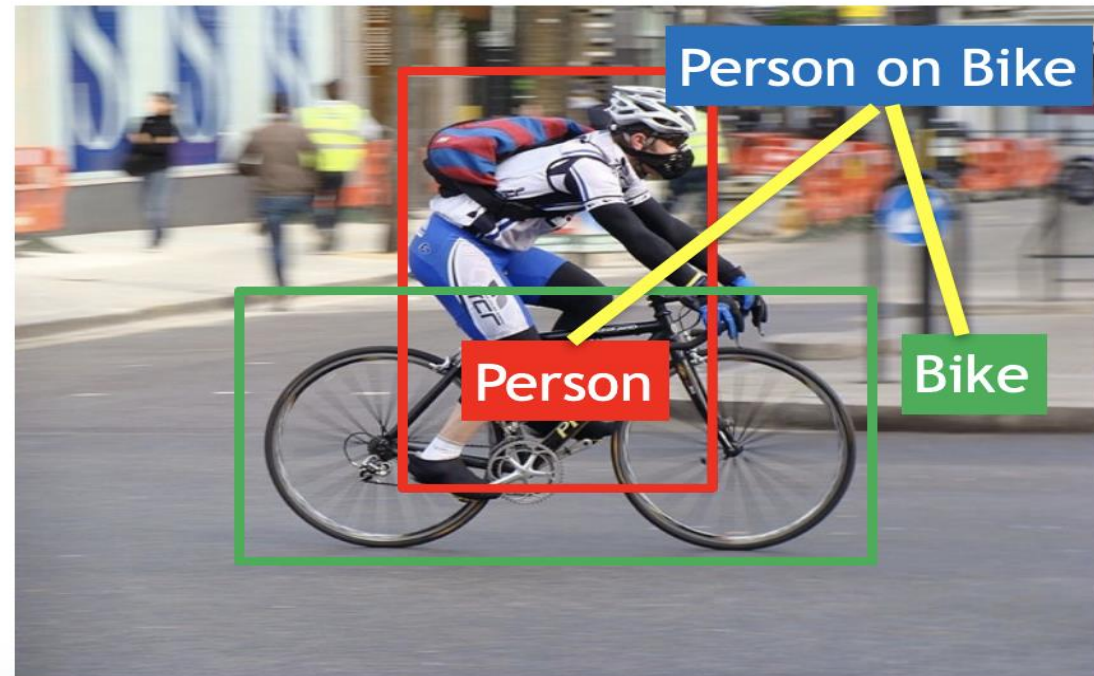
**Announcement:** Tomorrow's office hour is at 9am. Please let me know if you'd like to attend

# Coursework 1

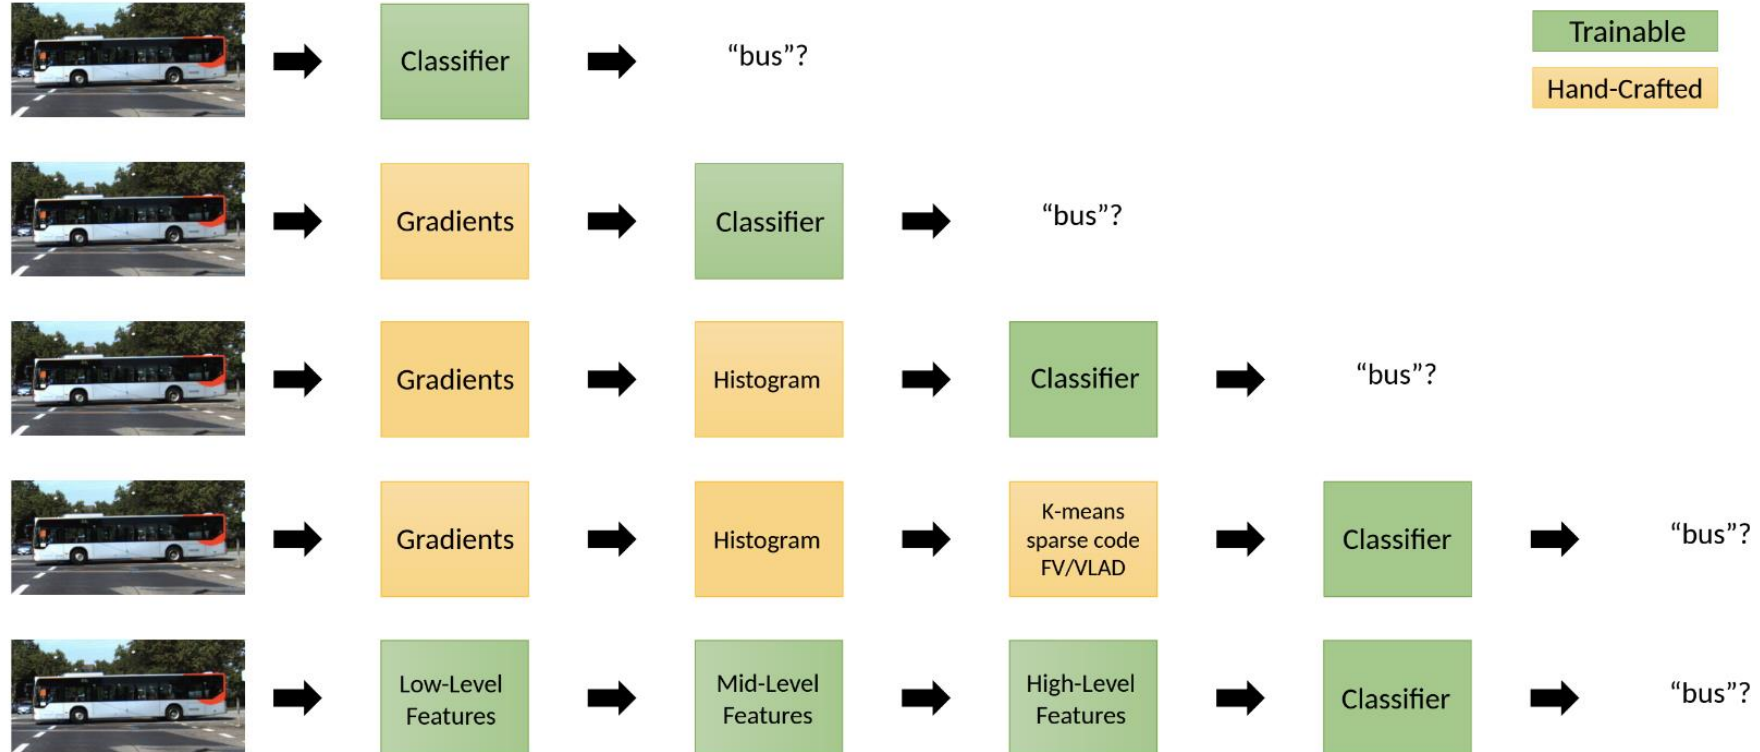
- First coursework (of two) released this Friday 18/10/2024
- Due 07/11/2024 at 16:00
- Accounts for 30% of the total mark for this module
- You will **only** be assessed on material covered in the **first 3 weeks** of lectures (i.e., including this lecture)
- Please refer to the UCL guidance on plagiarism and use of LLMs – these apply to this work!
- If you have any queries (at all!) please speak to me or a member of staff that you feel comfortable with

# Invariance and equivariance in computer vision

# Computer vision

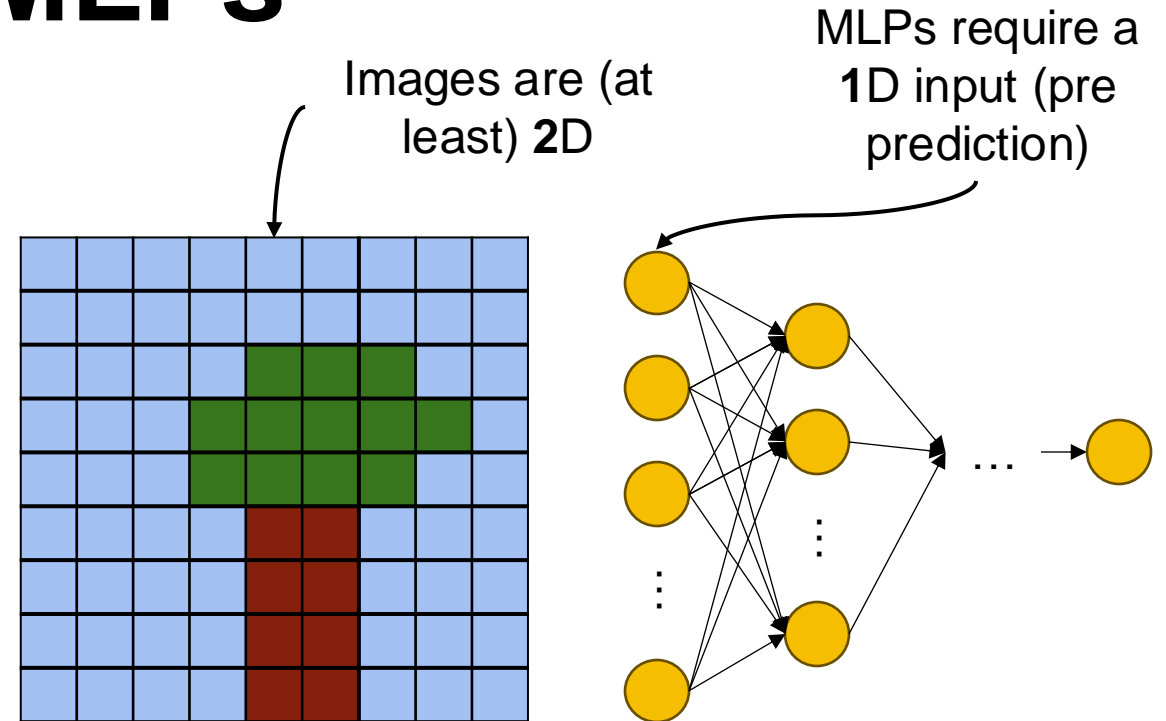


# Computer vision: Representation learning

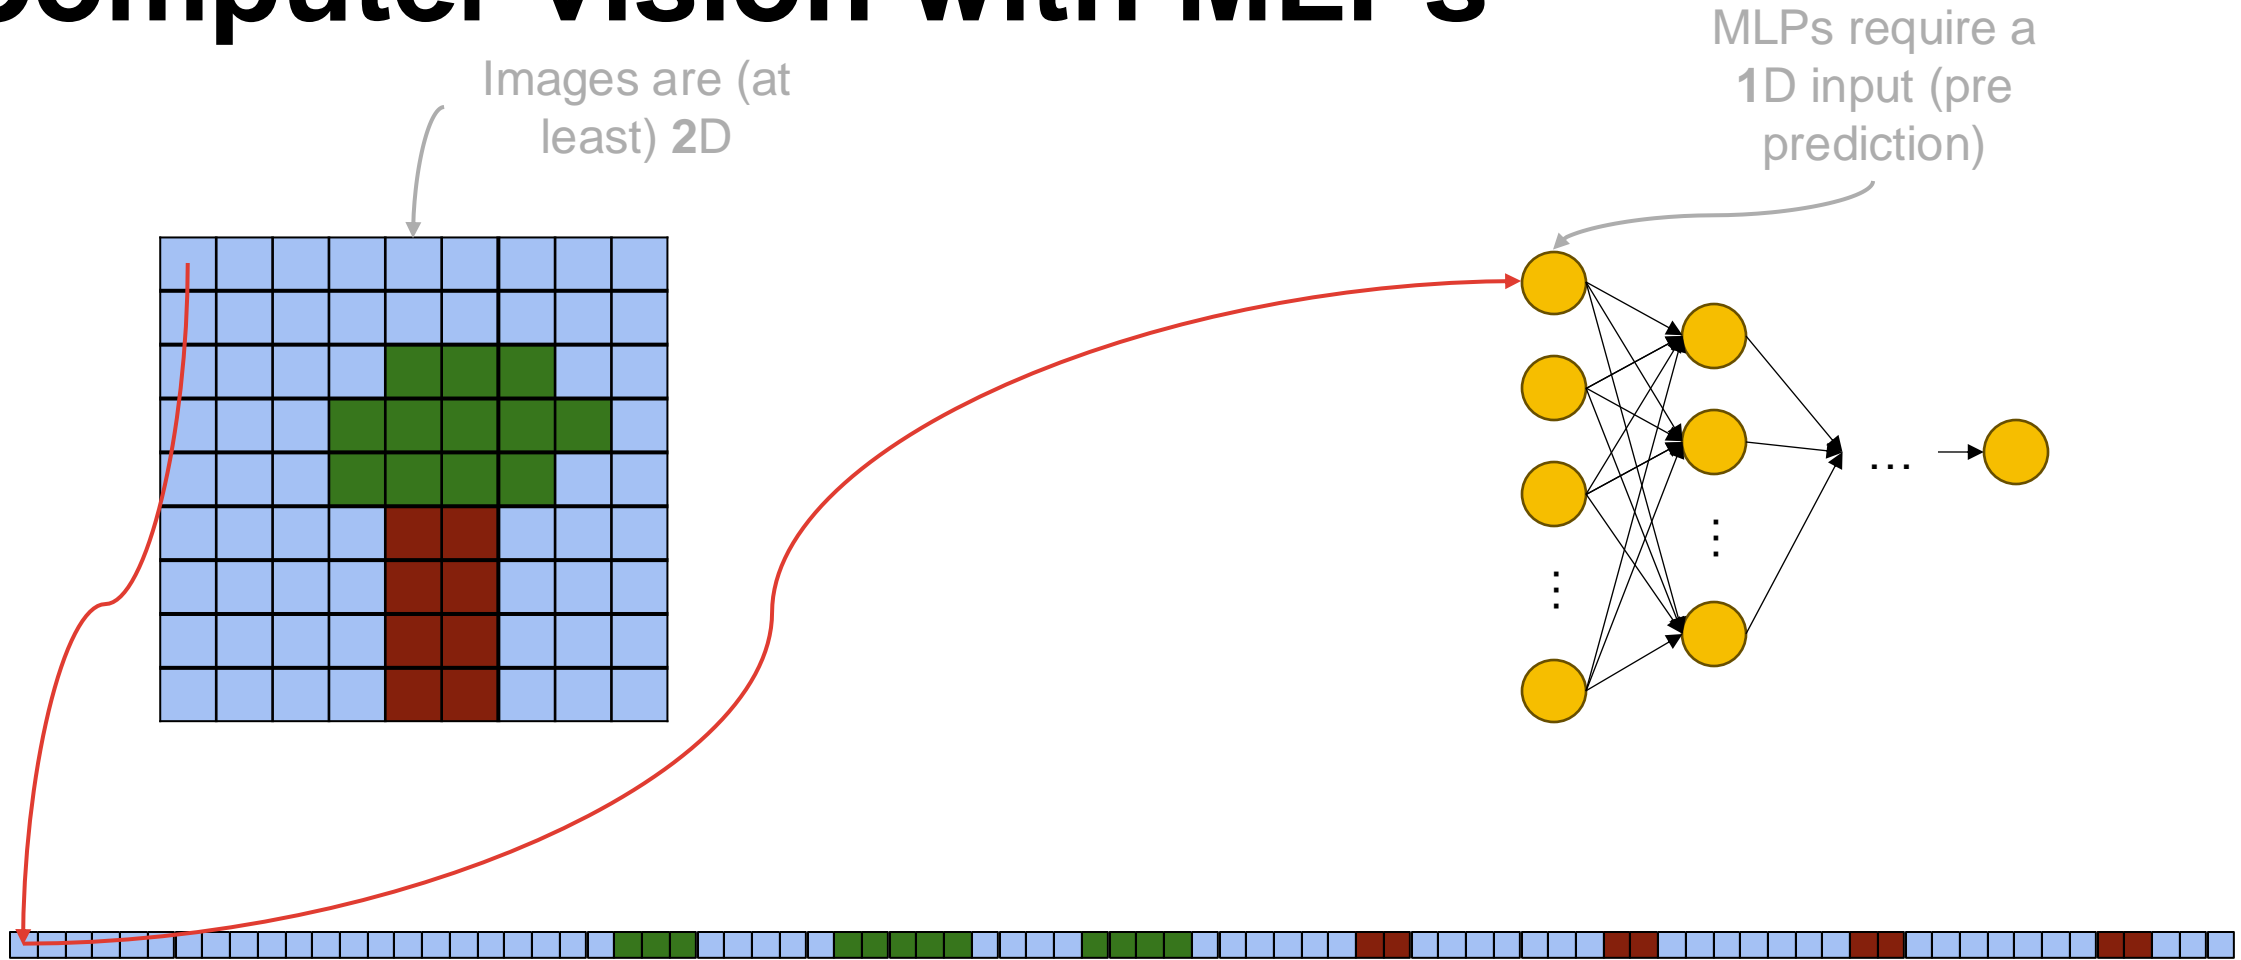


# Computer vision with MLPs

- For a given prediction, vanilla MLPs assume a  $d$ -dimensional input i.e.,  $d$  features
  - In the first lecture, we predicted the probability of you passing given 2 features
- Images are a  $d \times d$ -dimensional input



# Computer vision with MLPs

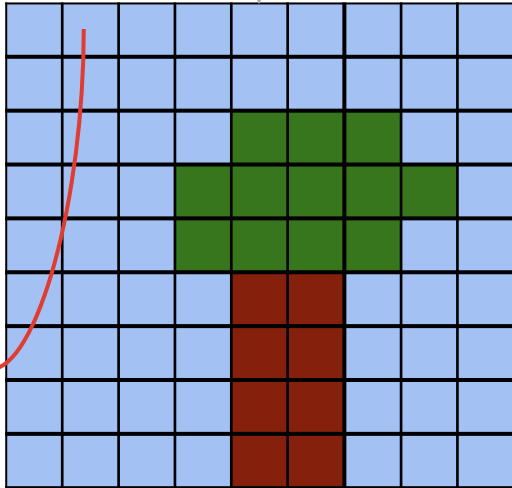


Flatten?

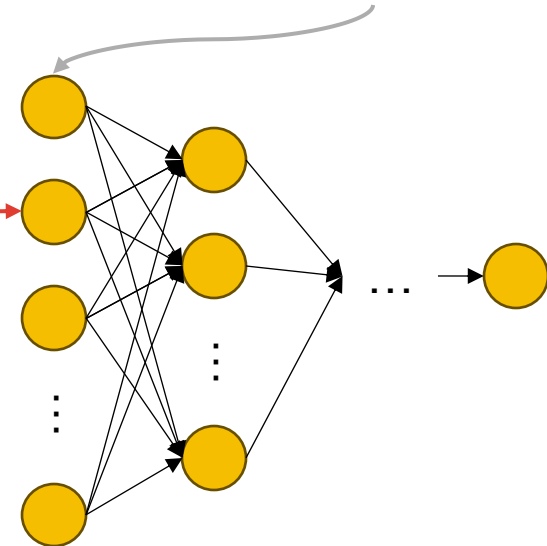


# Computer vision with MLPs

Images are (at least) 2D

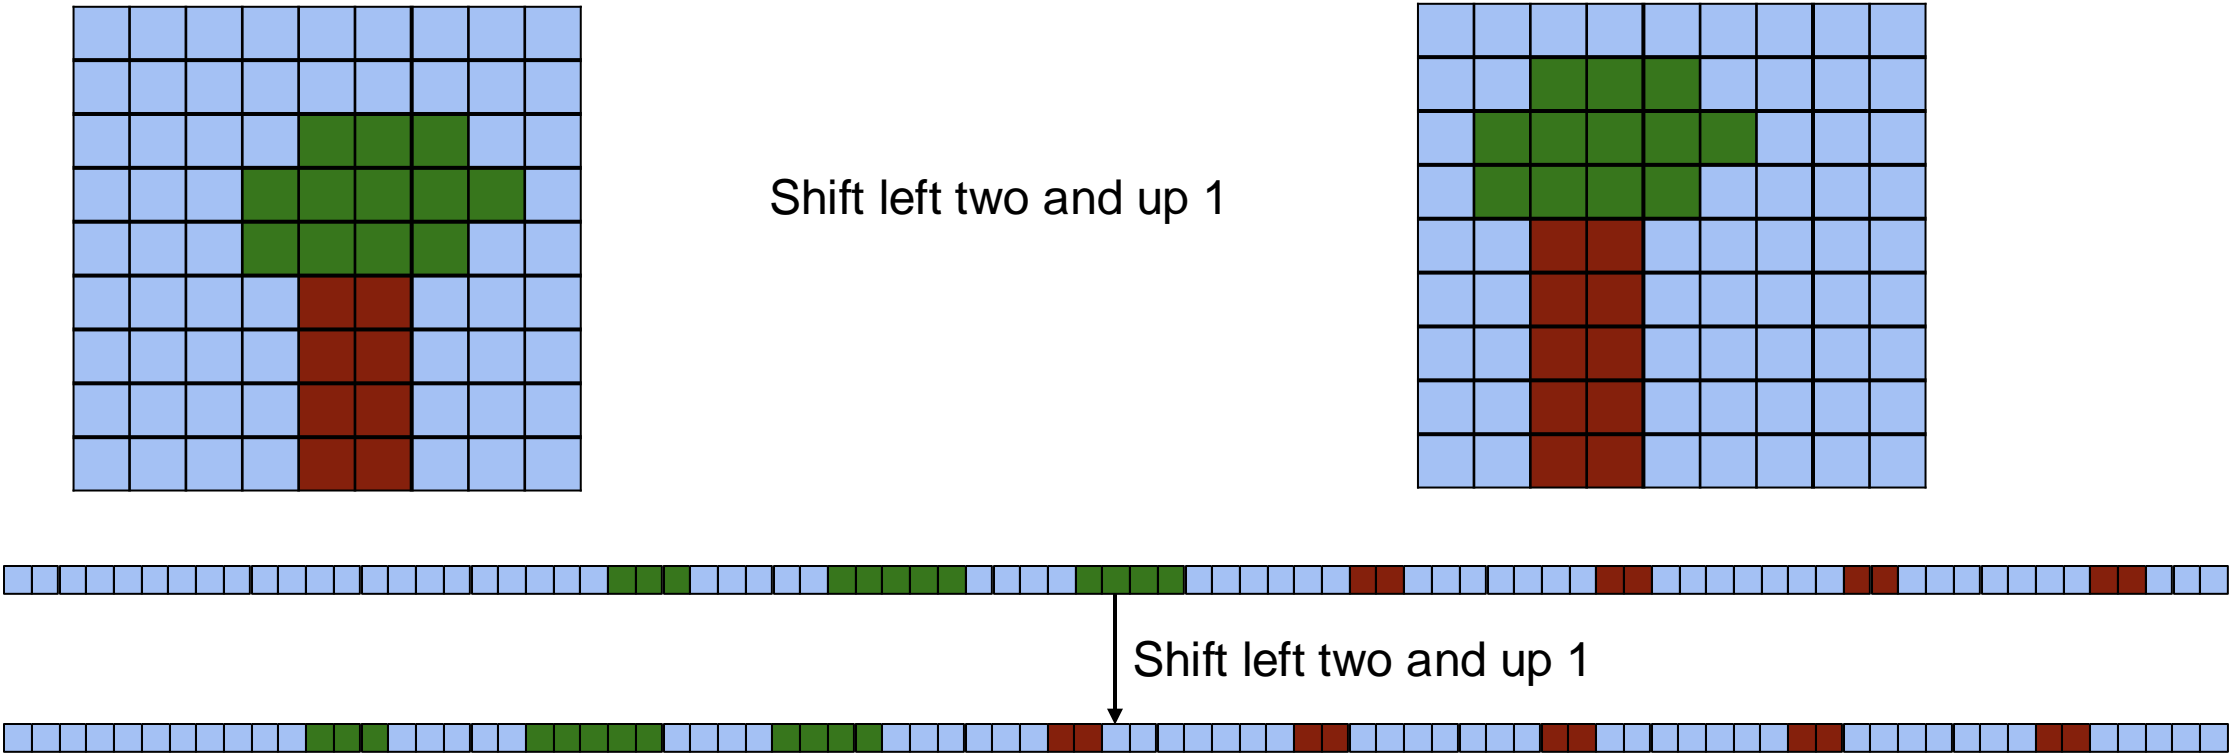


MLPs require a 1D input (pre prediction)



Flatten?

# Computer vision with MLPs

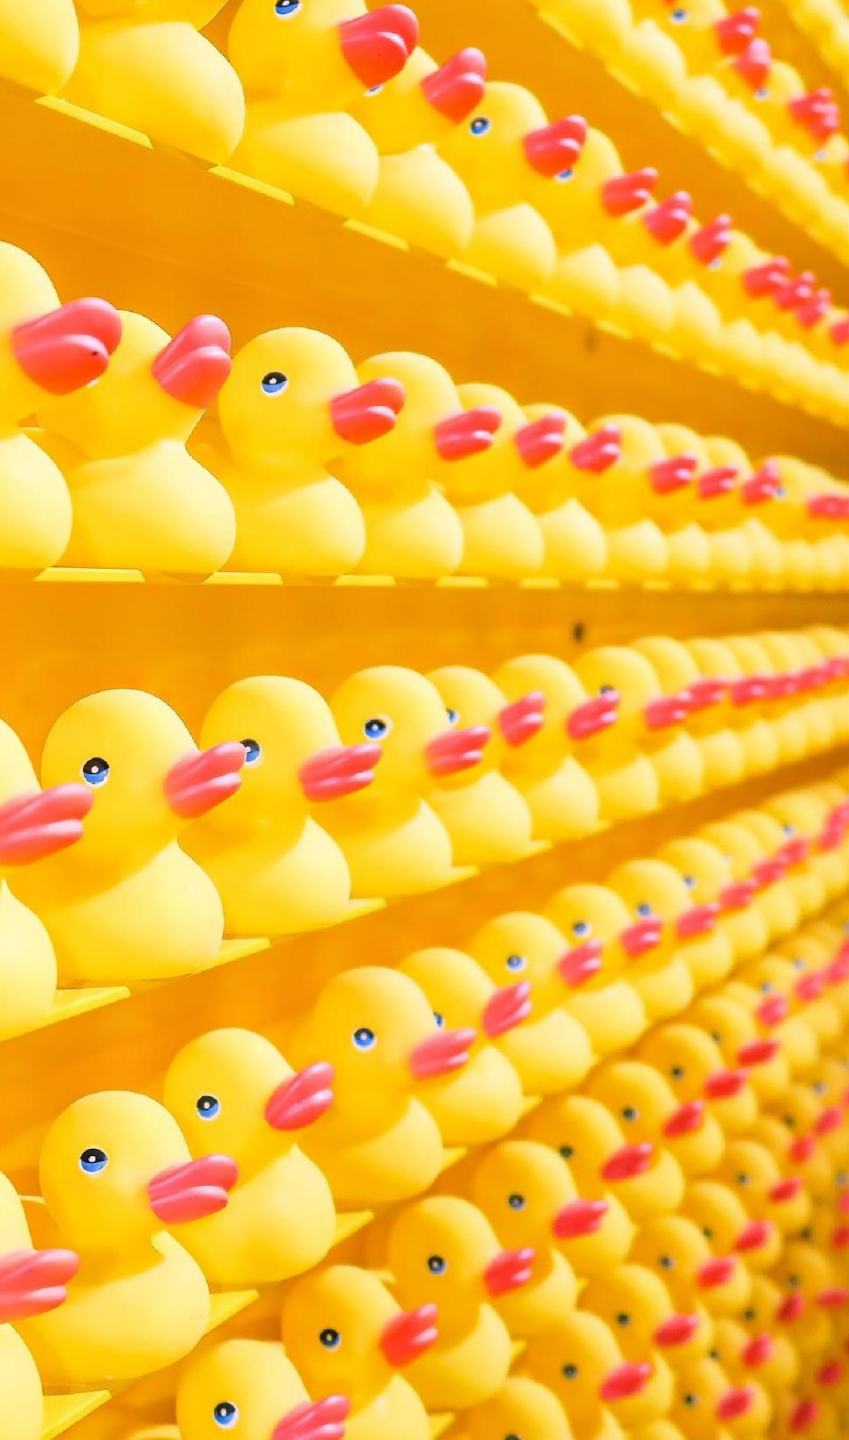


**Problem:** The model would need to see **every object** in **every position**

# Inductive biases

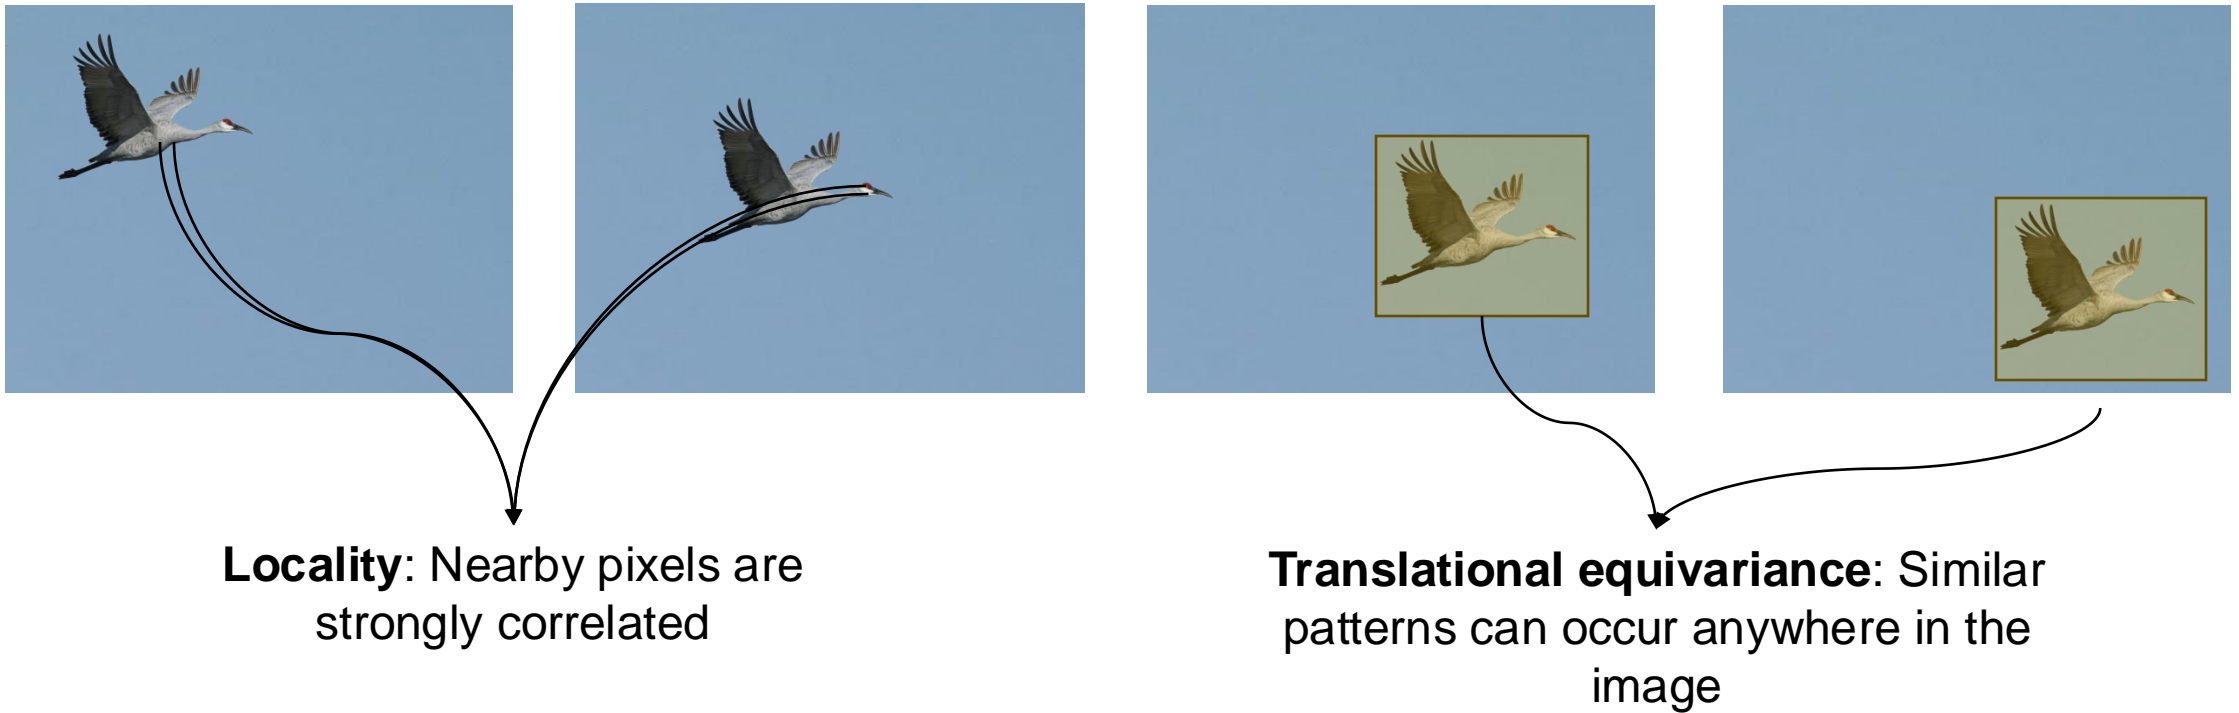
- Elements of model development
  - Function  $f: \mathcal{X} \rightarrow \mathcal{Y}$ , where  $f \in \mathcal{F}$ 
    - A (data)set of samples which have been taken from our input and output sets  $d = \{(x, y)_i: x \in \mathcal{X}, y \in \mathcal{Y}, i \in 1, \dots, n, n \in \mathbb{N}\}$ 
      - $a: \mathcal{X} \times \mathcal{Y} \times l \rightarrow \mathcal{F}$
      - $l: \hat{\mathcal{Y}} \times \mathcal{Y} \rightarrow \mathbb{R}$
- Each design choice imposes an **inductive bias** on the learning process
- Informally, inductive biases **impose a preference over  $\mathcal{F}$ /directly alter the shape of  $\mathcal{F}$**
- When the inductive bias **prefers a subset of  $\mathcal{F}$**  which contains a function  $\hat{f}$  that **obtains minimal generalisation error**, the inductive bias is **useful** and **improves learning**







# Utilising topological structure: Locality and translational equivariance



# Locality

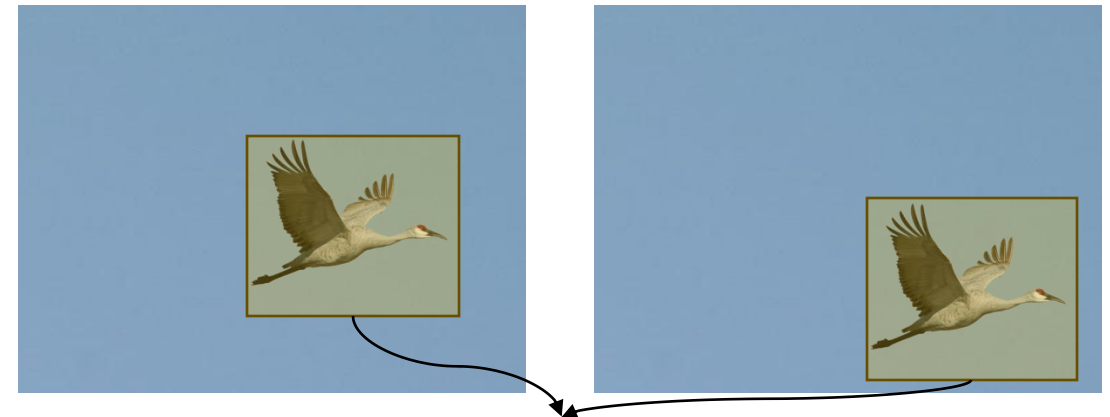
- A condition being local loosely means that its effect is restricted to objects in a neighbourhood
  - The importance of locality in image recognition is loosely that: we care about objects **within the image** not the entire image
- Formally defined in: Geometric Deep Learning Grids, Groups, Graphs, Geodesics, and Gauges, 2021



**Locality:** Nearby pixels are strongly correlated

# Translational equivariance

- **Invariance**: the output of function  $f$  is **unaffected** by a “transformation” of the input
  - $f(\rho x) = f(x)$
- **Equivariance**: The input and output are **affected in the same way** by a transformation on the input
  - $f(\rho x) = \rho f(x)$
- References:
  - Geometric Deep Learning Grids, Groups, Graphs, Geodesics, and Gauges, 2021
  - Understanding deep learning, 2024

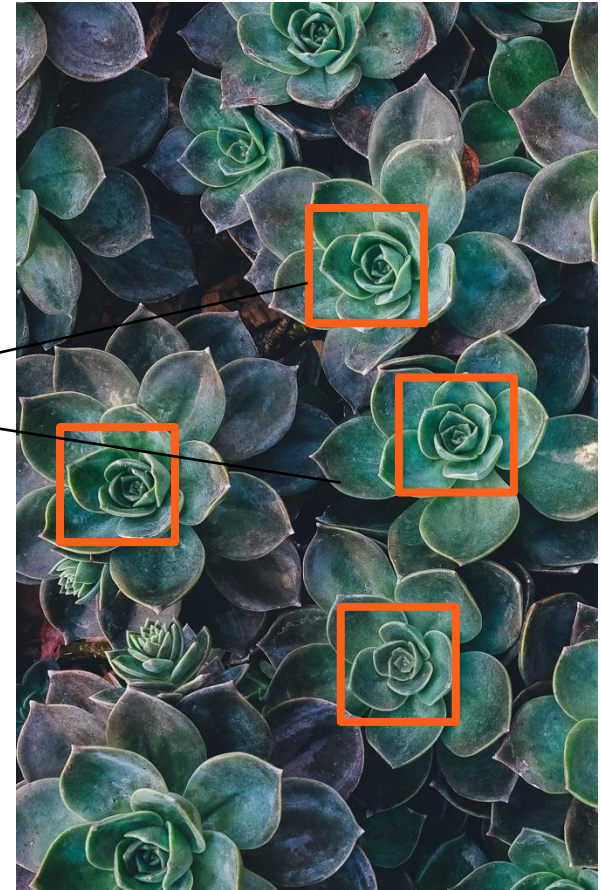


**(Local) translational invariance/equivariance:**  
Similar patterns can occur anywhere in the image

- Identifying whether an image contains a bird: require translational **invariance**
- Identify bounding box of bird: require translational **equivariance**

# Implementing (local) translational equivariance

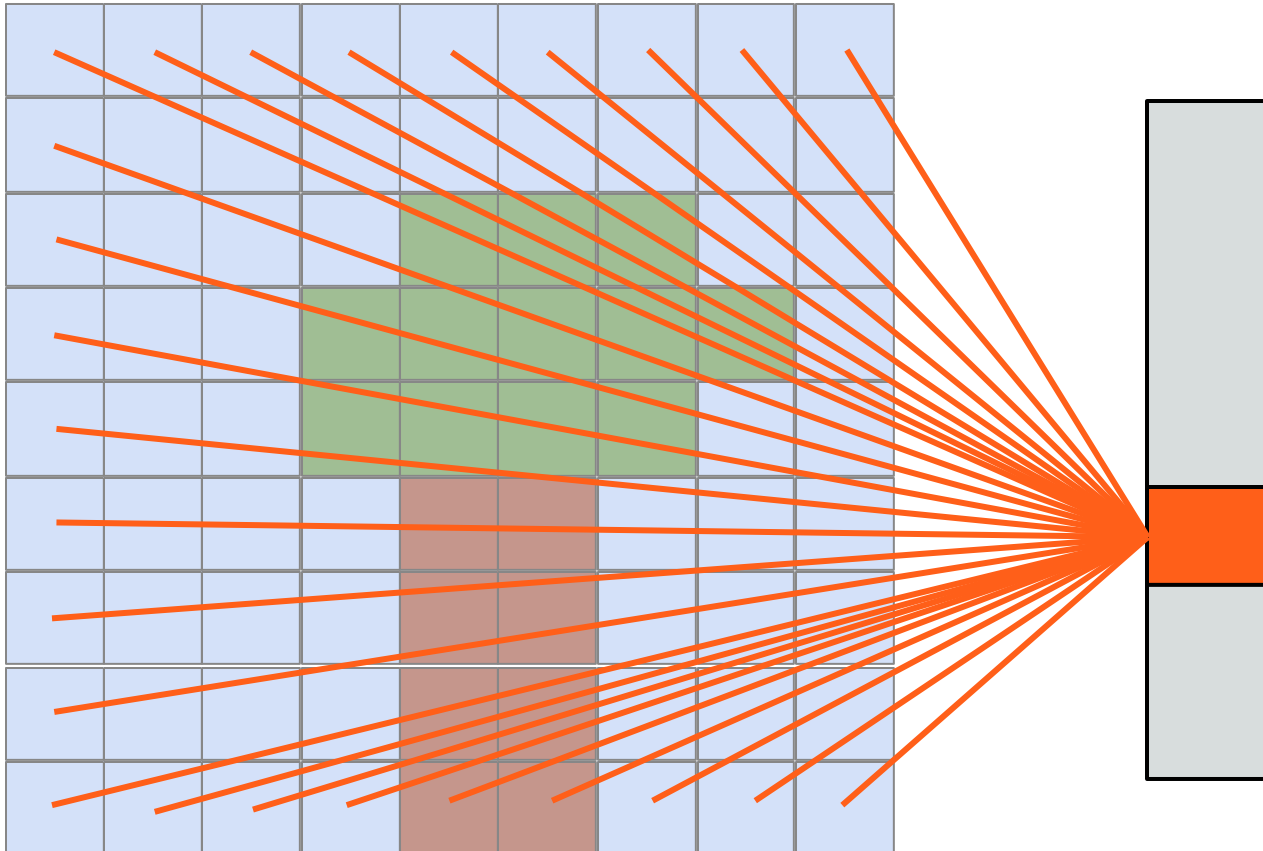
- **Weight sharing**
  - Neural network weights are **activated** in the same way **if the same input is provided**
- Applying the **same weights** to these blocks of input will result in the **same values for the activations**





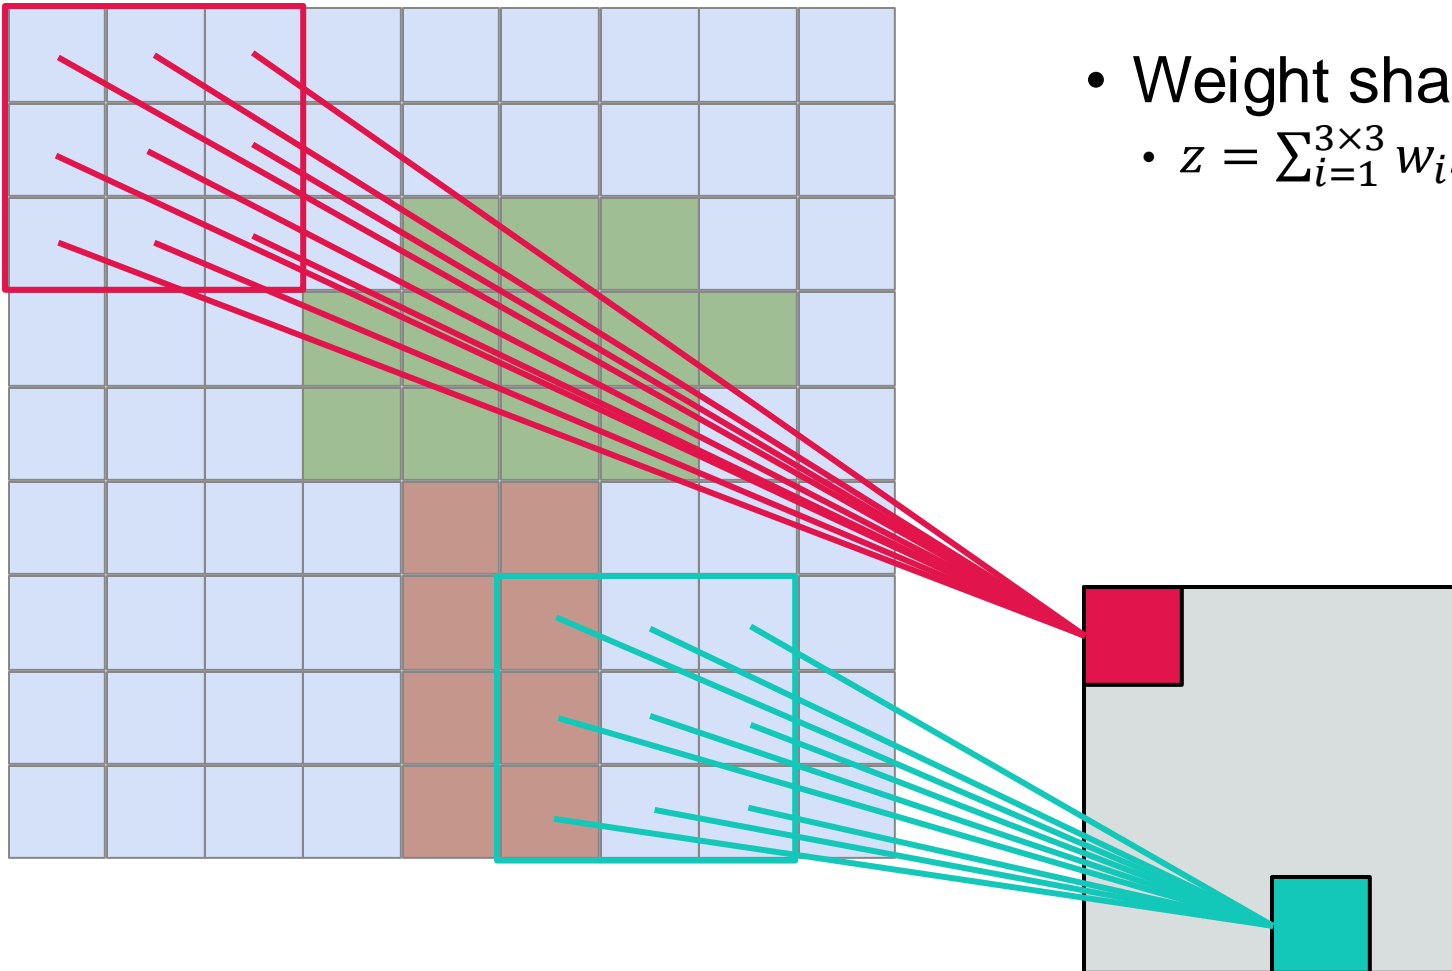
# CNNs

# From fully connected to locally connected: weight sharing



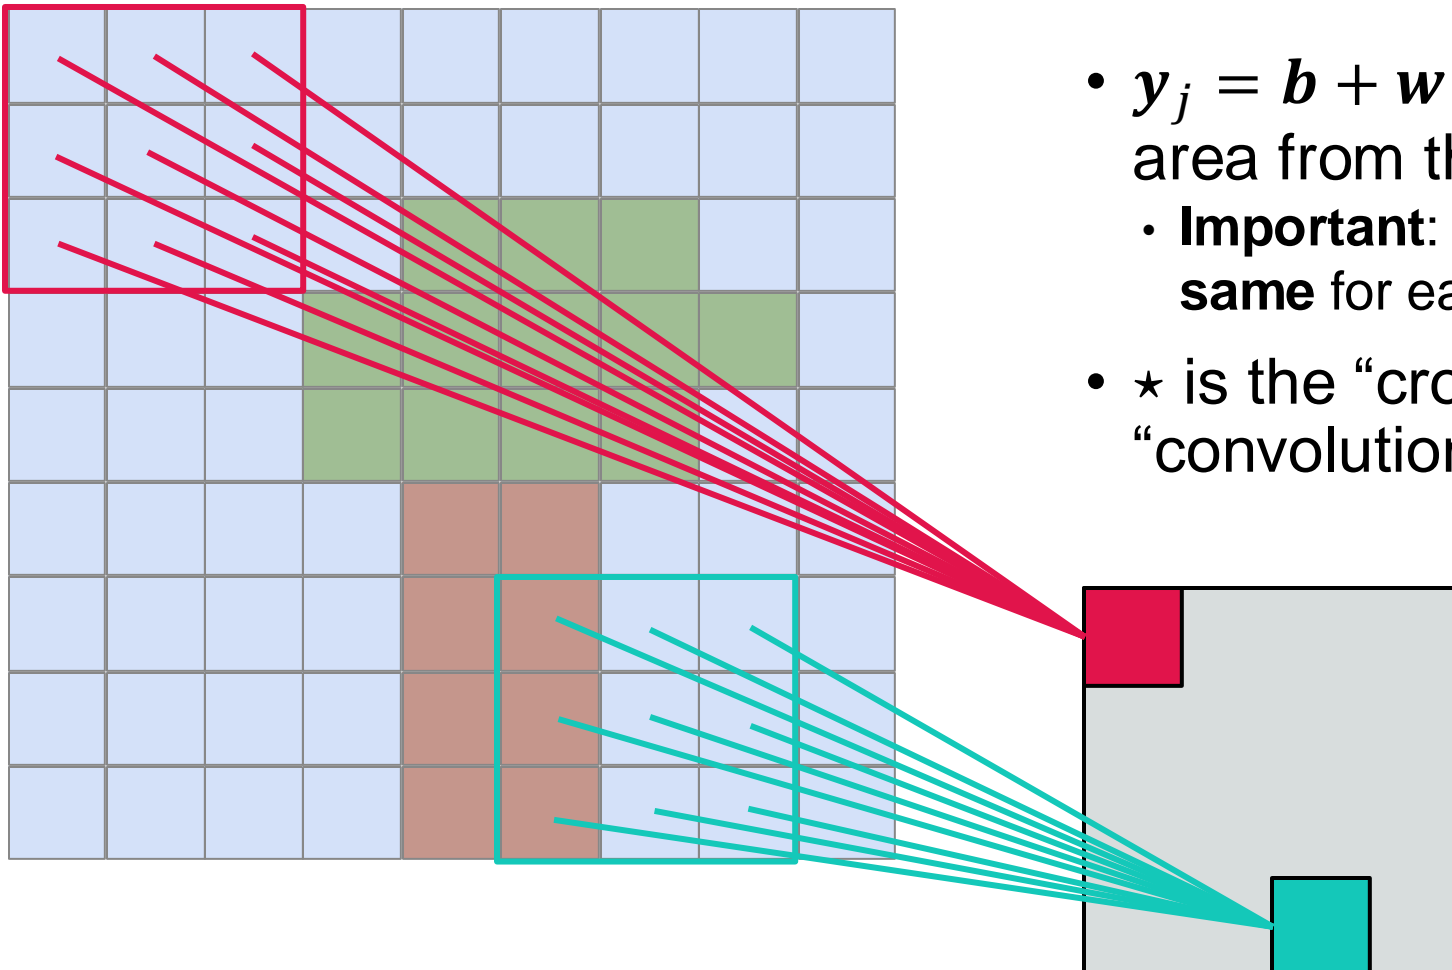
- Fully connected unit
  - $z = \sum_{i=1}^{|d| \times |d|} w_i x_i + b$

# From fully connected to locally connected: weight sharing



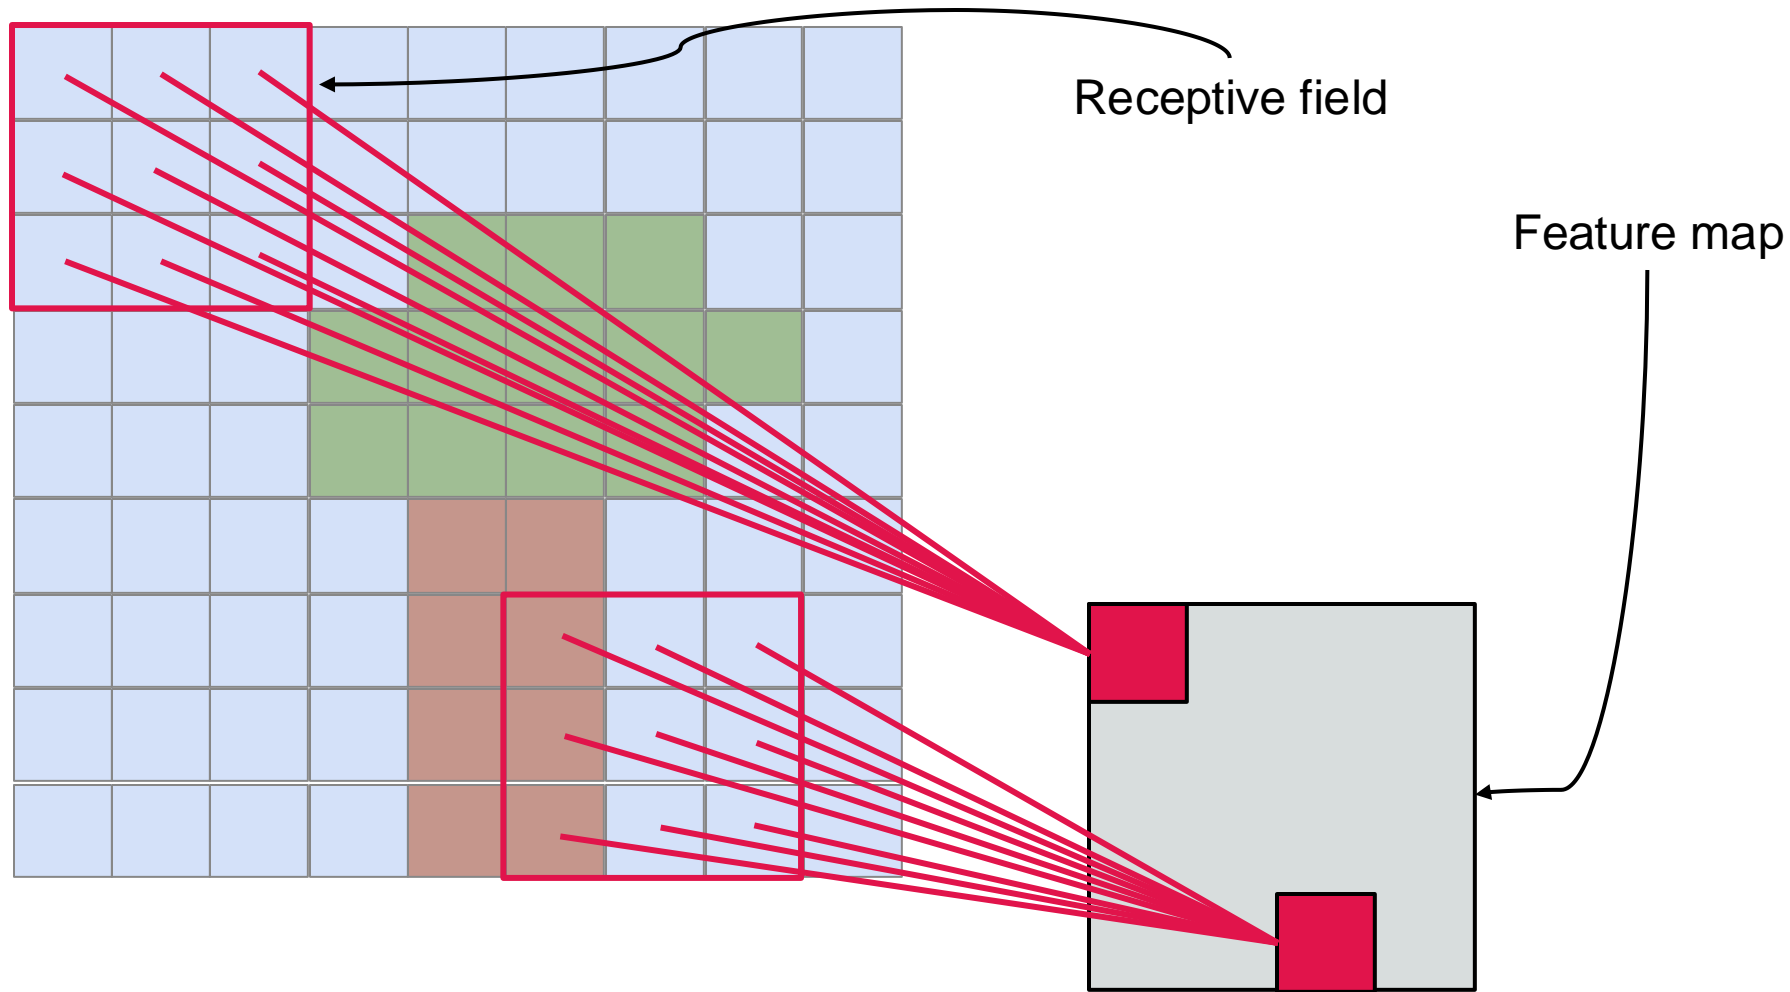
- Weight sharing over  $3 \times 3$  filter
  - $z = \sum_{i=1}^{3 \times 3} w_i x_i + b$

# Convolution

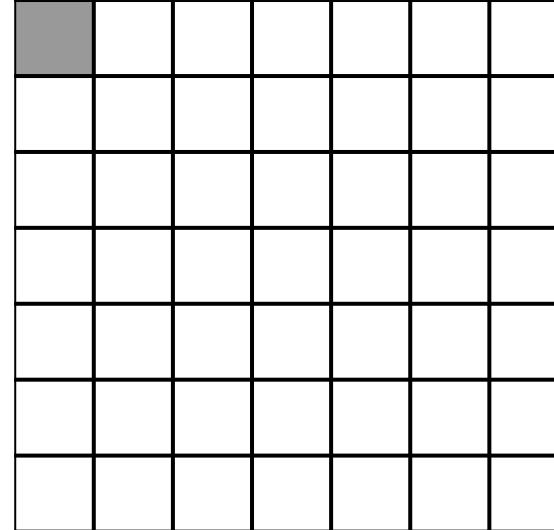
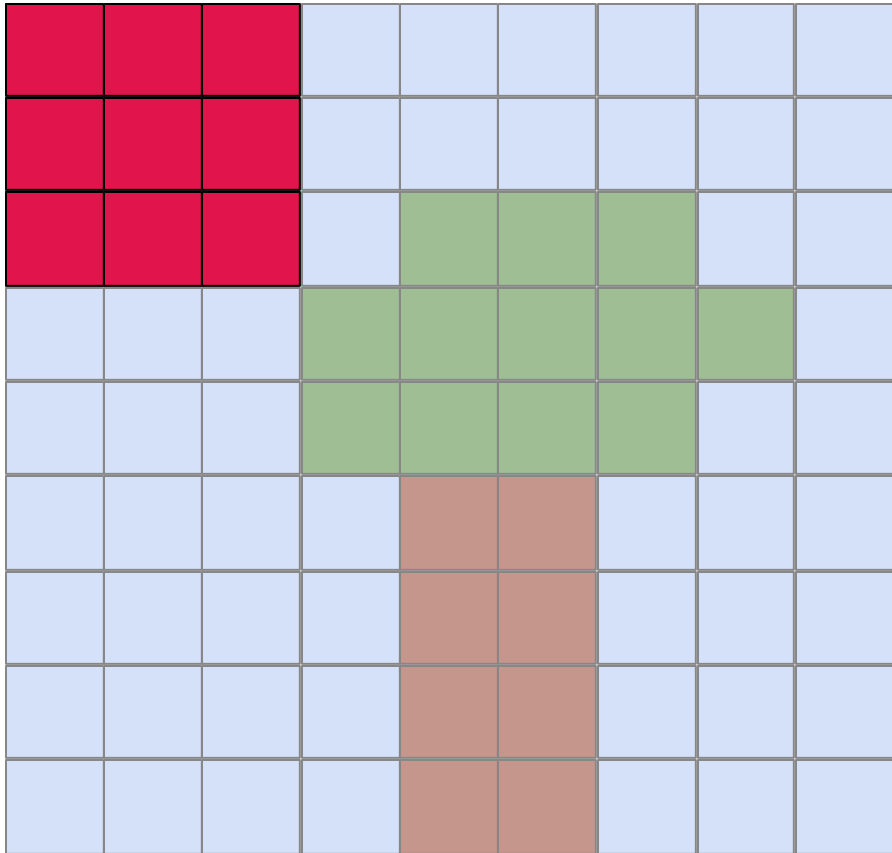


- $y_j = b + w \star x_j$  where  $x_j$  is a local  $3 \times 3$  area from the input image
  - **Important:**  $b$  and  $w$  (the weights) are the **same** for each  $3 \times 3$  input area
- $\star$  is the “cross-correlation” but called “convolution” in ML
- Extended exercise:
  - Global cross-correlation is defined as  $(f \star g)(\tau) = \int_{-\infty}^{\infty} f(t)g(t + \tau)dt$
  - Obtain the discrete local version from the previous slide

# Convolution

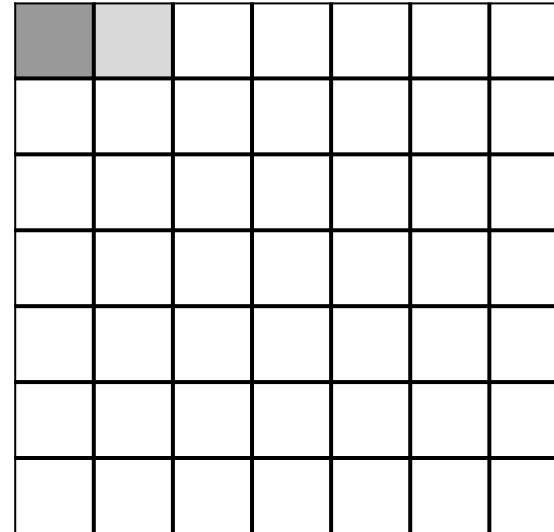
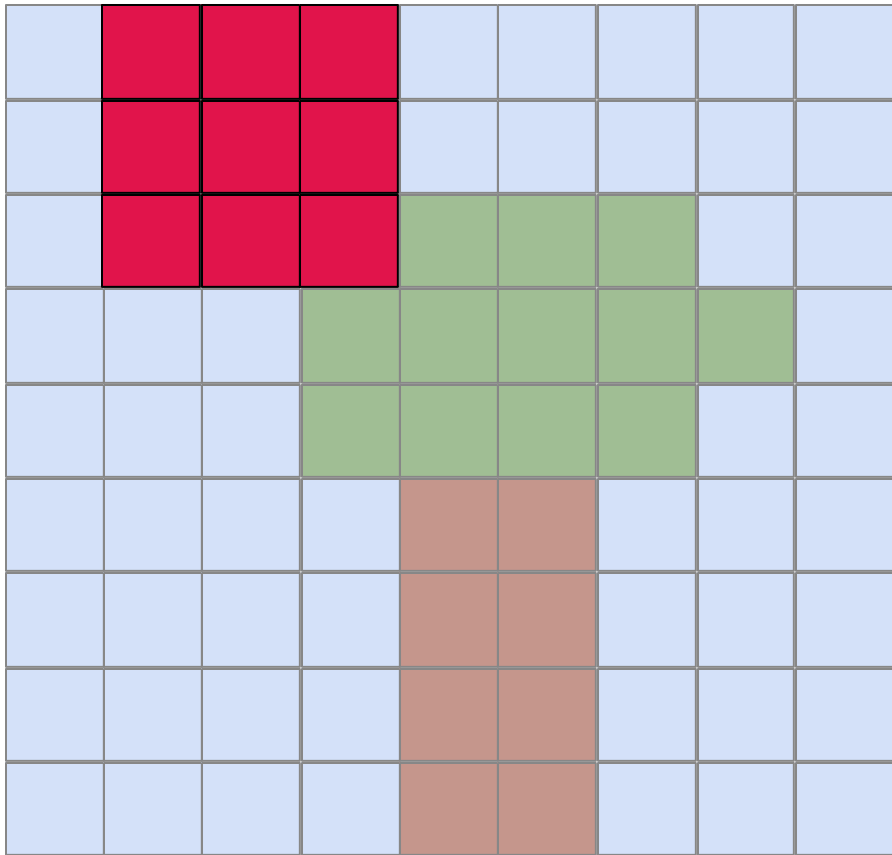


# Implementing the convolution



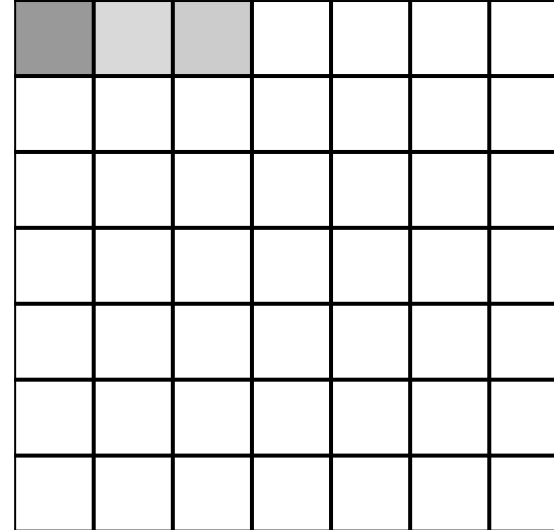
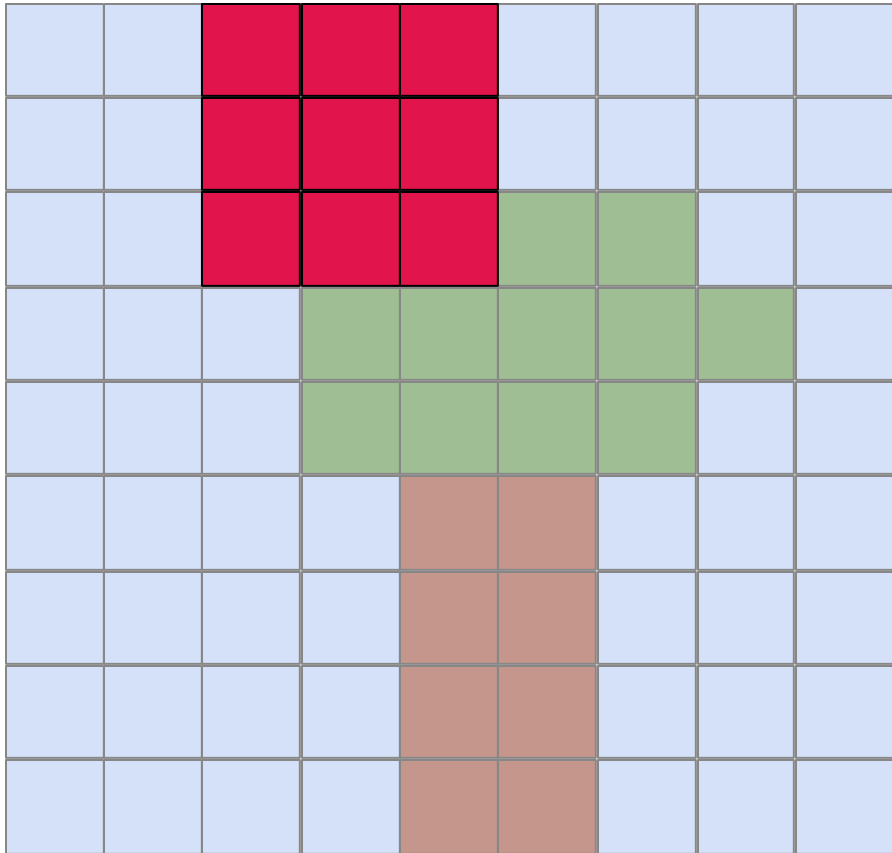
The **kernel/filter** slides across the image and produces an output value for each position

# Implementing the convolution



The **kernel/filter** slides across the image and produces an output value for each position

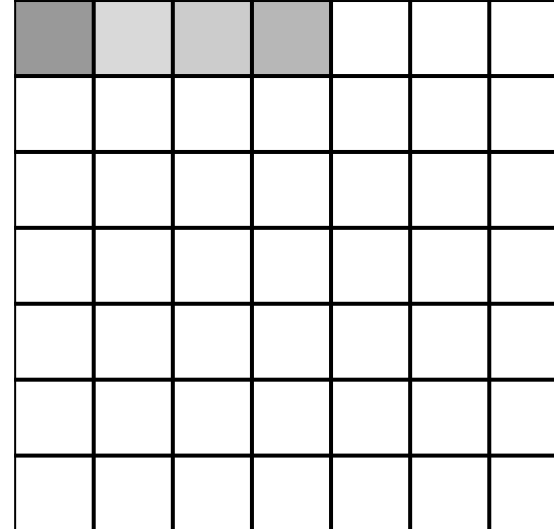
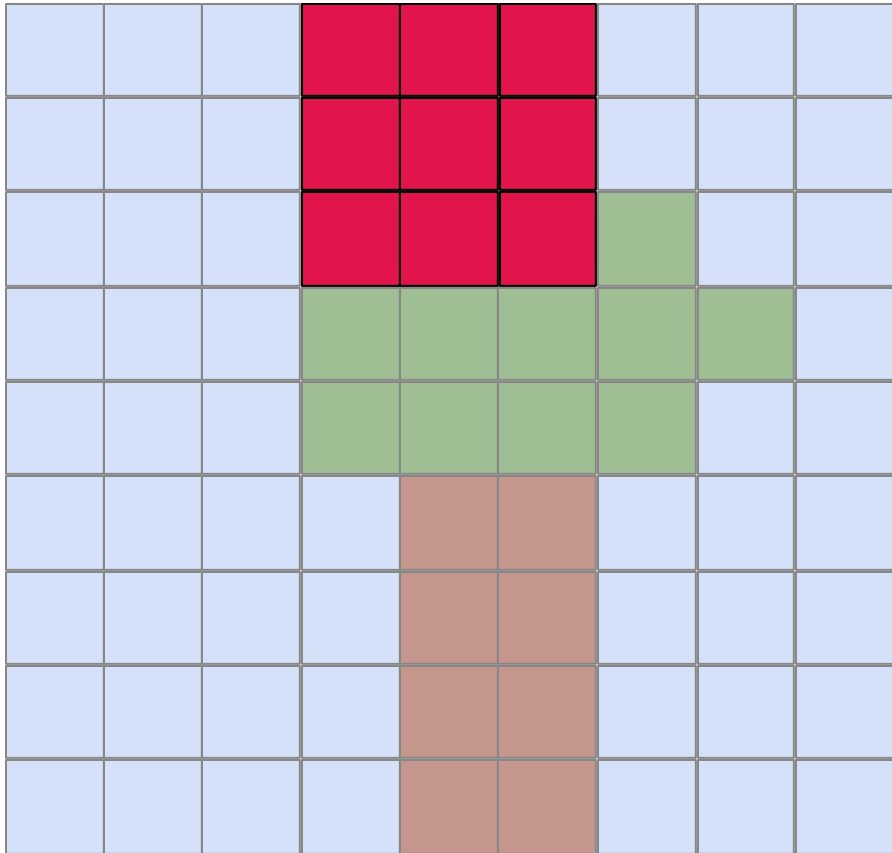
# Implementing the convolution



The **kernel/filter** slides across the image and produces an output value for each position

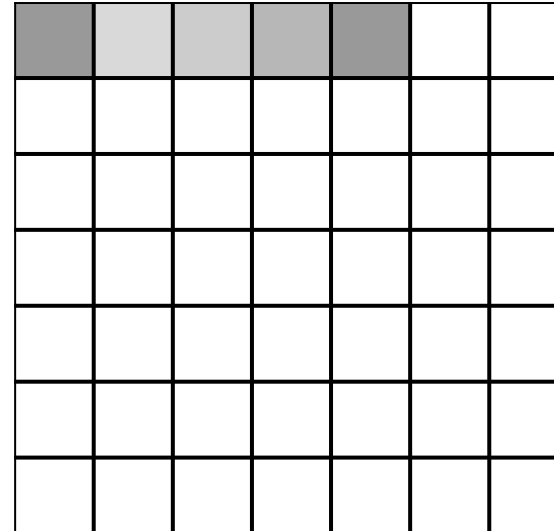
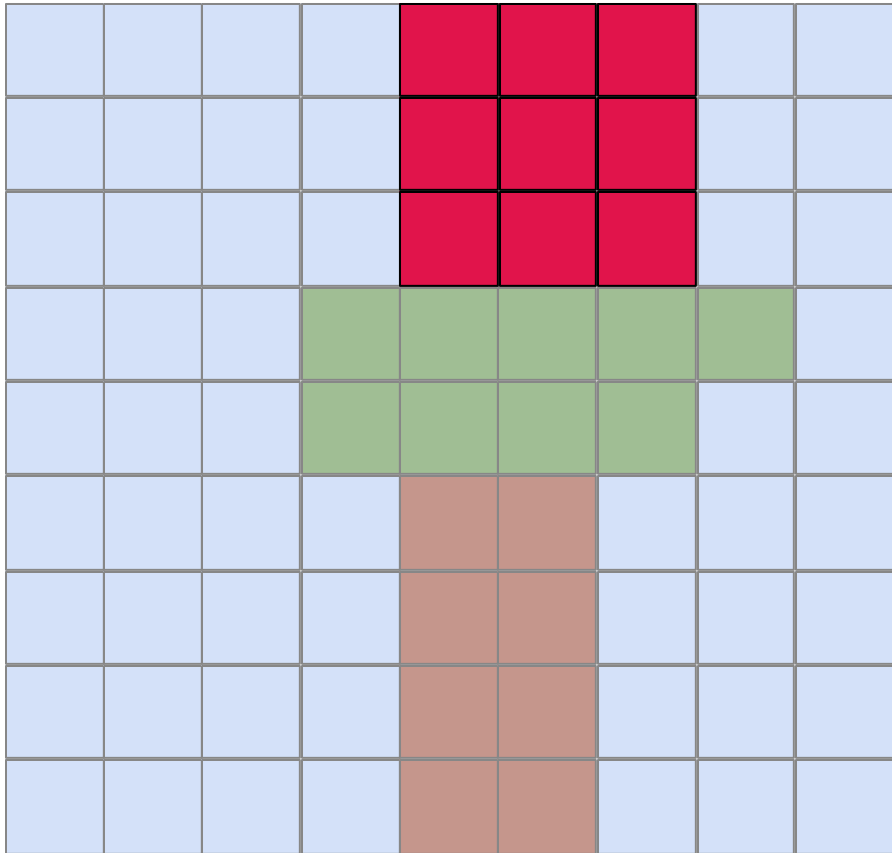


# Implementing the convolution



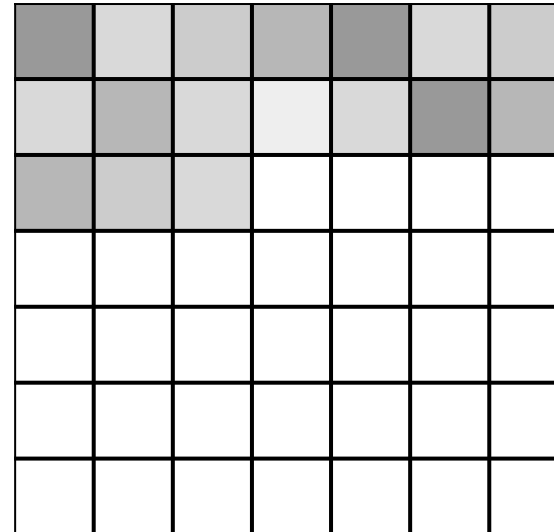
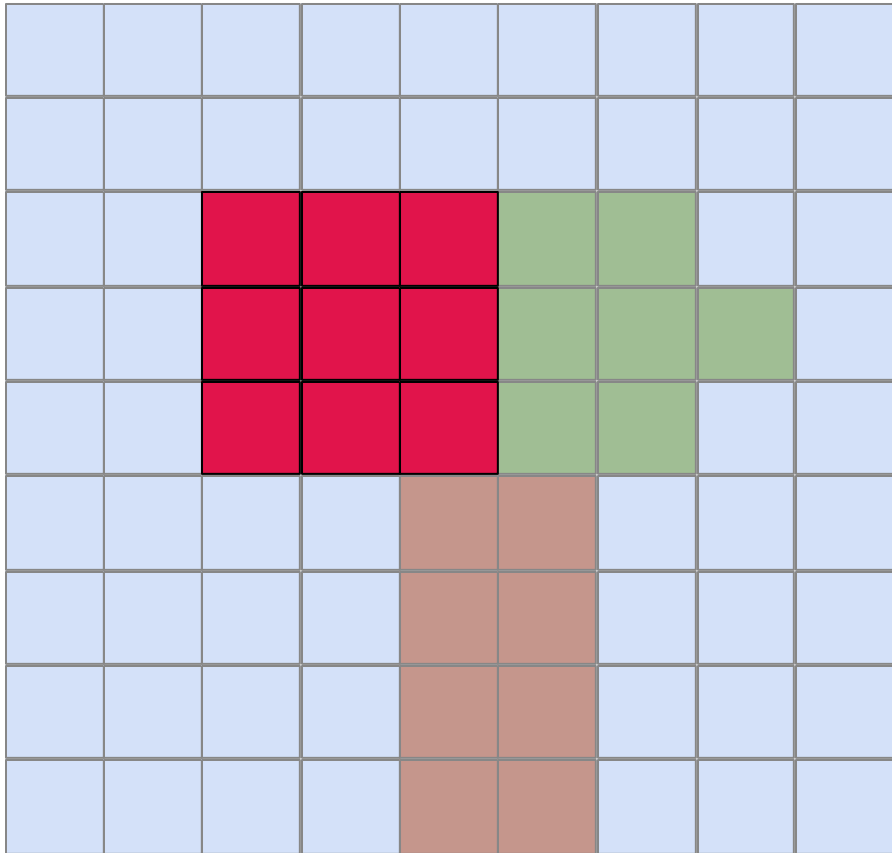
The **kernel/filter** slides across the image and produces an output value for each position

# Implementing the convolution



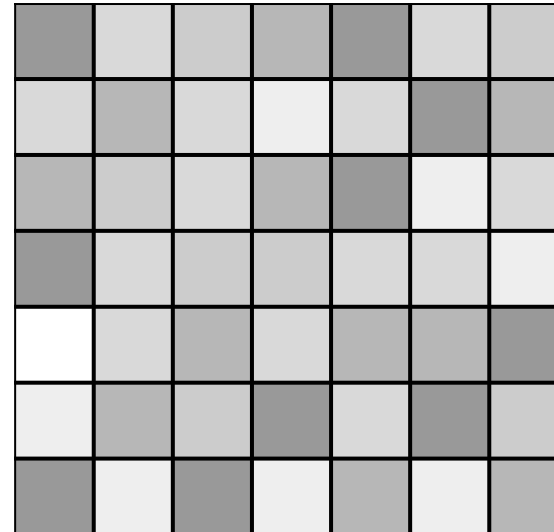
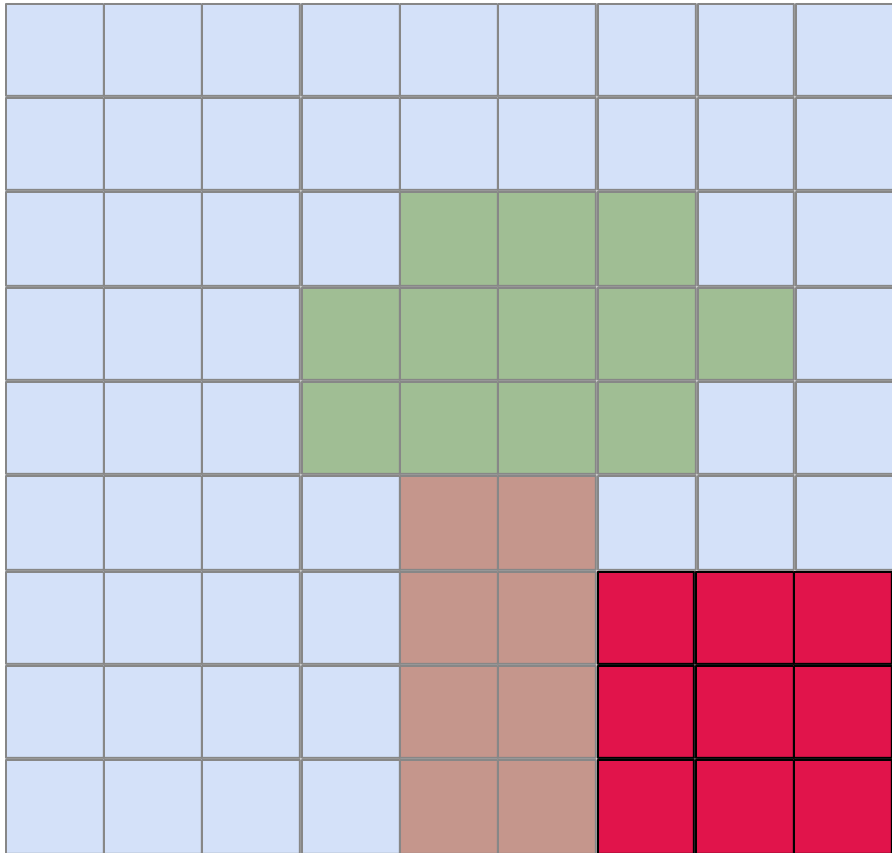
The **kernel/filter** slides across the image and produces an output value for each position

# Implementing the convolution



The **kernel/filter** slides across the image and produces an output value for each position

# Implementing the convolution



The **kernel/filter** slides across the image and produces an output value for each position

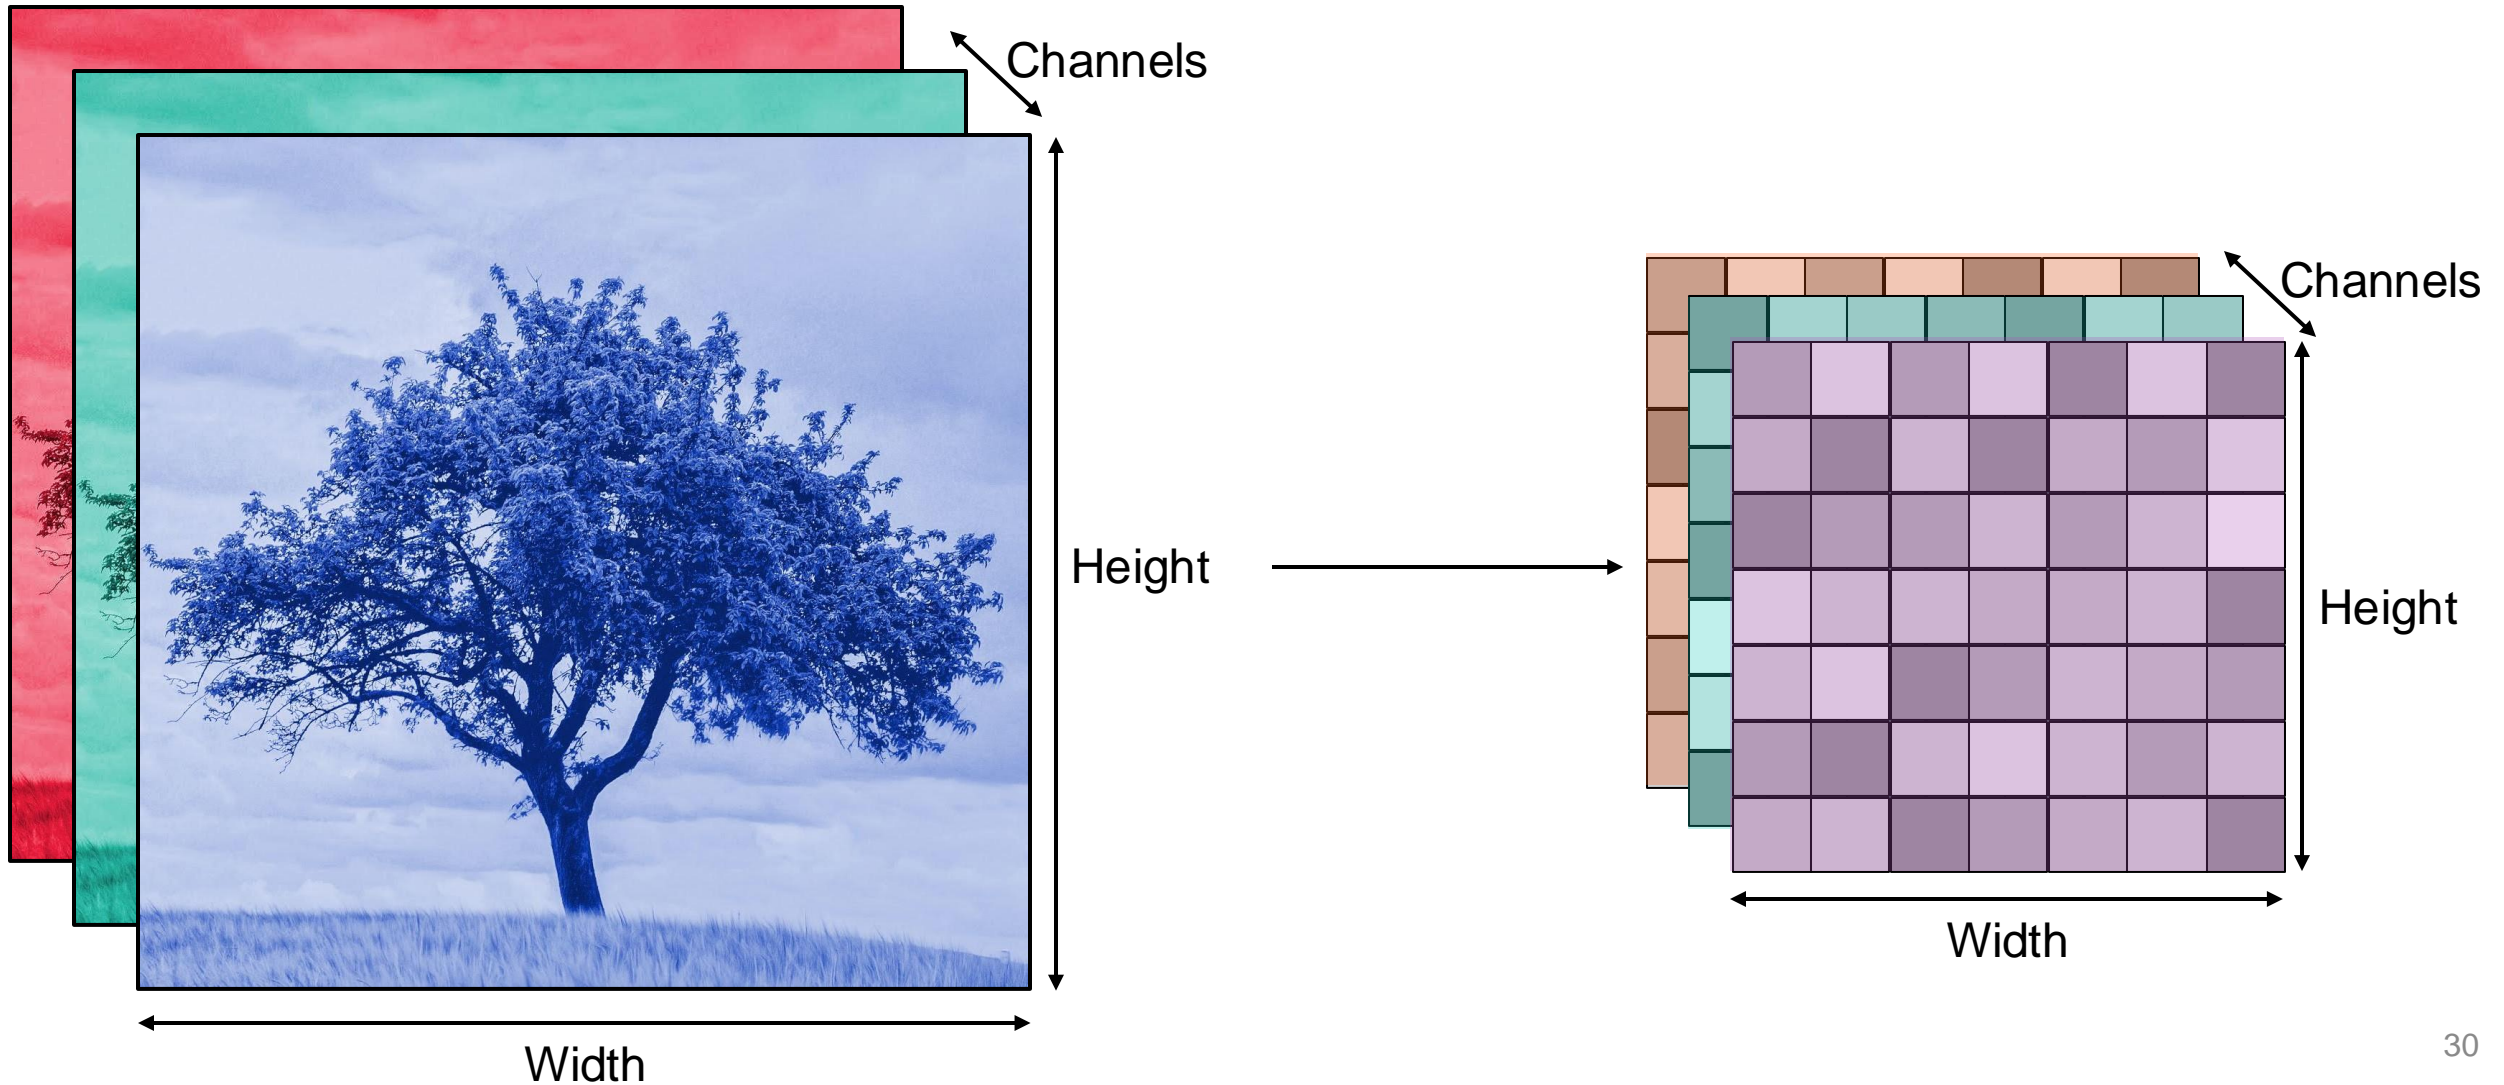
# Pictures as tensors



RGB



# Pictures as tensors





# CNN in pytorch

```
class CNN(nn.Module):

    def __init__(self):
        super().__init__()
        module_list = nn.ModuleList()
        module_list.append(nn.Conv2d(
            in_channels=1,
            out_channels=1,
            kernel_size=(3, 3)
        ))

        module_list.append(nn.Flatten())

        module_list.append(
            nn.Linear(900,1)
        )
        module_list.append(
            nn.Sigmoid()
        )
        self.module_list = module_list

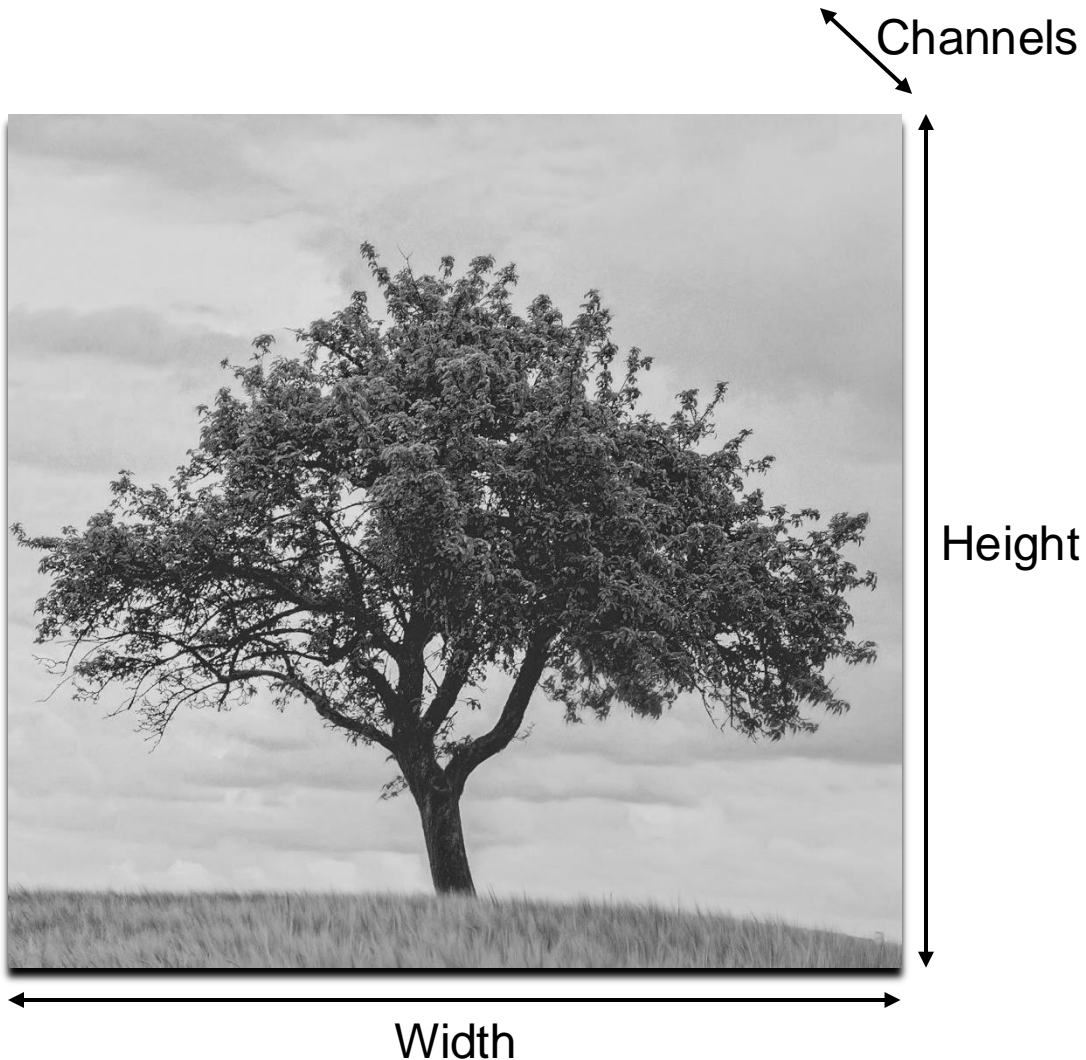
    def forward(self,x:torch.Tensor)->torch.Tensor:
        _x = x
        for l in self.module_list:
            _x = l(_x)
        return _x
```

```
model = CNN()
summary(model, (1,1,32,32))

✓ 0.0s
```

Layer (type:depth-idx)	Output Shape	Param #
CNN	[1, 1]	--
└ModuleList: 1-1	--	--
└└Conv2d: 2-1	[1, 1, 30, 30]	10
└└Flatten: 2-2	[1, 900]	--
└└Linear: 2-3	[1, 1]	901
└└Sigmoid: 2-4	[1, 1]	--
Total params: 911		
Trainable params: 911		
Non-trainable params: 0		
Total mult-adds (Units.MEGABYTES): 0.01		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.01		
Params size (MB): 0.00		
Estimated Total Size (MB): 0.02		

# CNN in pytorch



```

model = CNN()
summary(model, (1,1,32,32))
✓ 0.0s

```

Layer (type:depth-idx)	Output Shape
CNN	[1, 1]
└ModuleList: 1-1	--
└└Conv2d: 2-1	[1, 1, 30, 30]
└└Flatten: 2-2	[1, 900]
└└Linear: 2-3	[1, 1]
└└Sigmoid: 2-4	[1, 1]

```

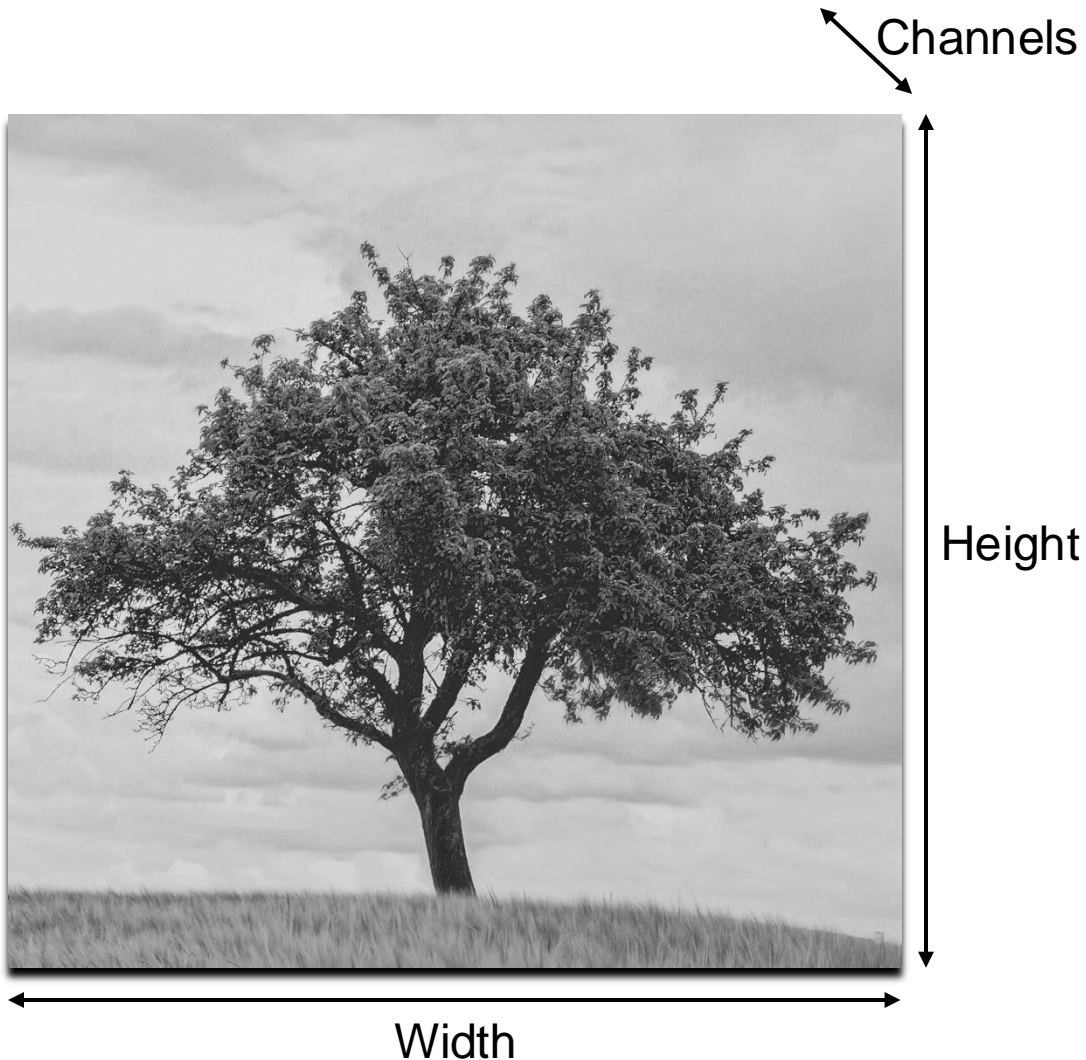
Total params: 911
Trainable params: 911
Non-trainable params: 0
Total mult-adds (Units.MEGABYTES): 0.01

Input size (MB): 0.00
Forward/backward pass size (MB): 0.01
Params size (MB): 0.00
Estimated Total Size (MB): 0.02

```



# CNN in pytorch



Batch size Channels Height Width

```
model = CNN()
summary(model, (1, 1, 32, 32))
```

✓ 0.0s

Layer (type:depth-idx)	Output Shape
CNN	[1, 1]
└ModuleList: 1-1	--
└└Conv2d: 2-1	[1, 1, 30, 30]
└└Flatten: 2-2	[1, 900]
└└Linear: 2-3	[1, 1]
└└Sigmoid: 2-4	[1, 1]

Total params: 911  
 Trainable params: 911  
 Non-trainable params: 0  
 Total mult-adds (Units.MEGABYTES): 0.01

Input size (MB): 0.00  
 Forward/backward pass size (MB): 0.01  
 Params size (MB): 0.00  
 Estimated Total Size (MB): 0.02

# Hierarchical local patterns

- **Hierarchical patterns**

- Local low-level features are composed into larger, more abstract features



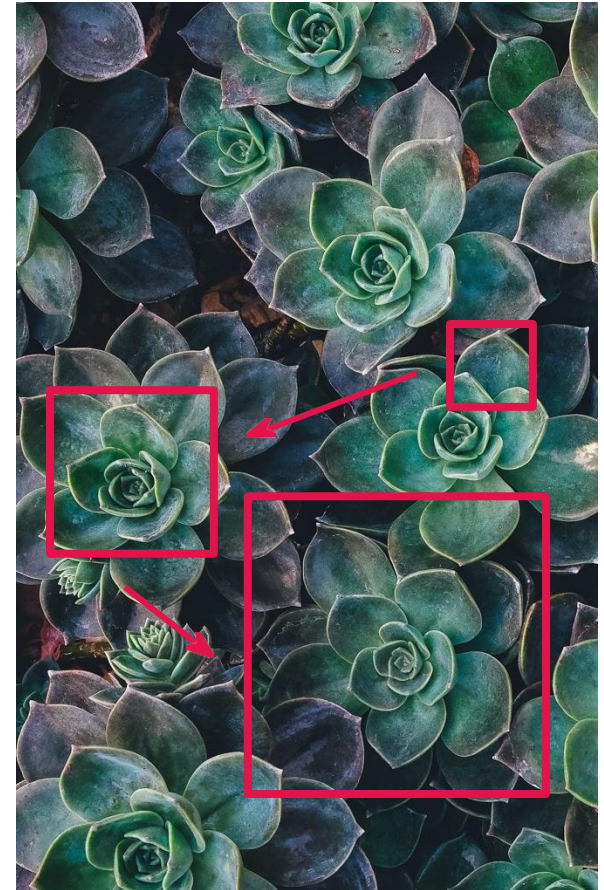
Edges and textures



Object parts

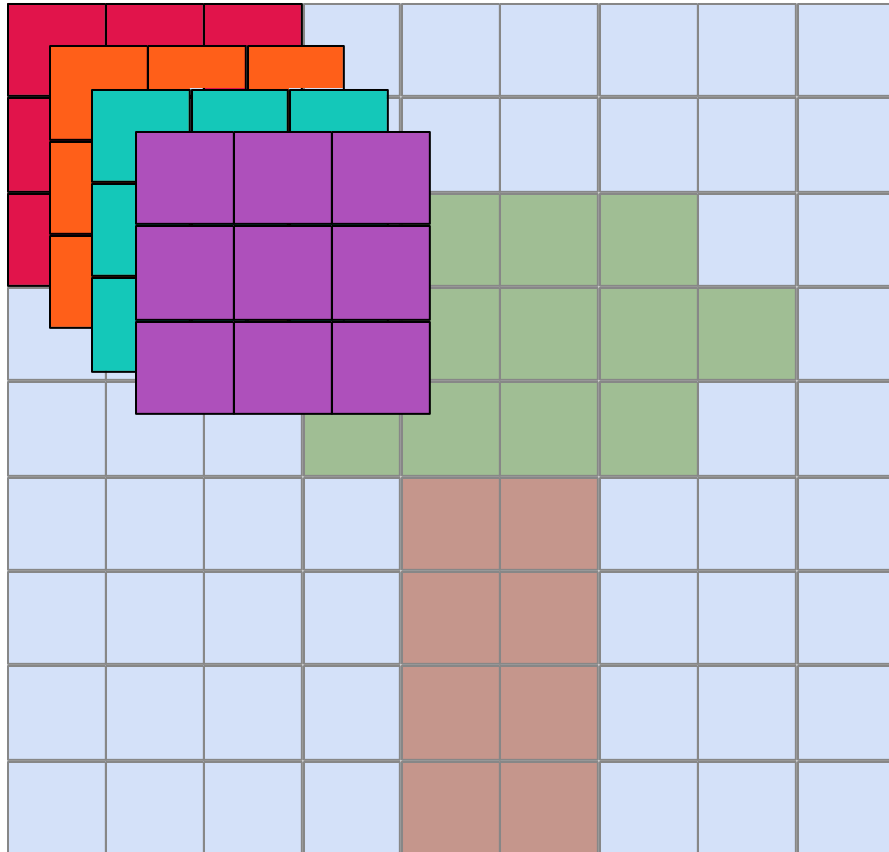


Objects

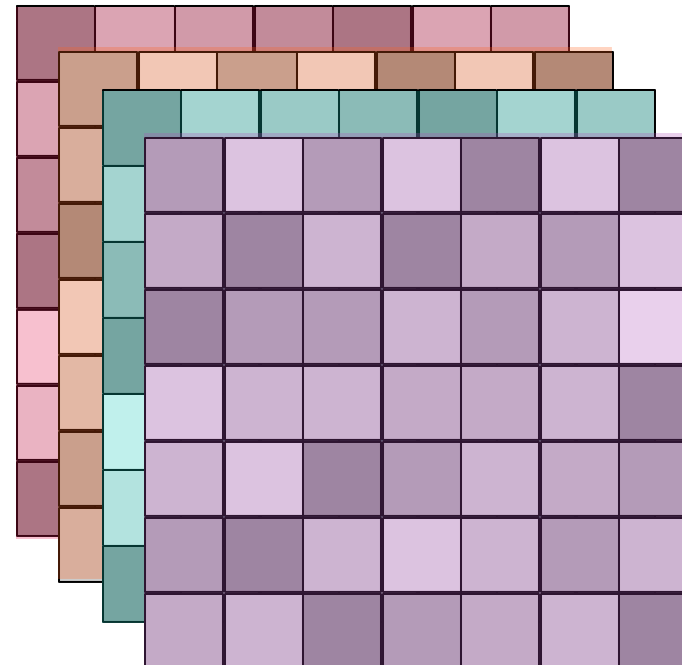


# Deeper and more complex feature maps

Input dimension: (1,1,9,9)



Output dimension: (1,4,9,9)



Different filters  
(weights) activate  
for different  
patterns

# Defining convolutional kernels

## Conv2d

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) \[SOURCE\]
```

Applies a 2D convolution over an input signal composed of several input planes.

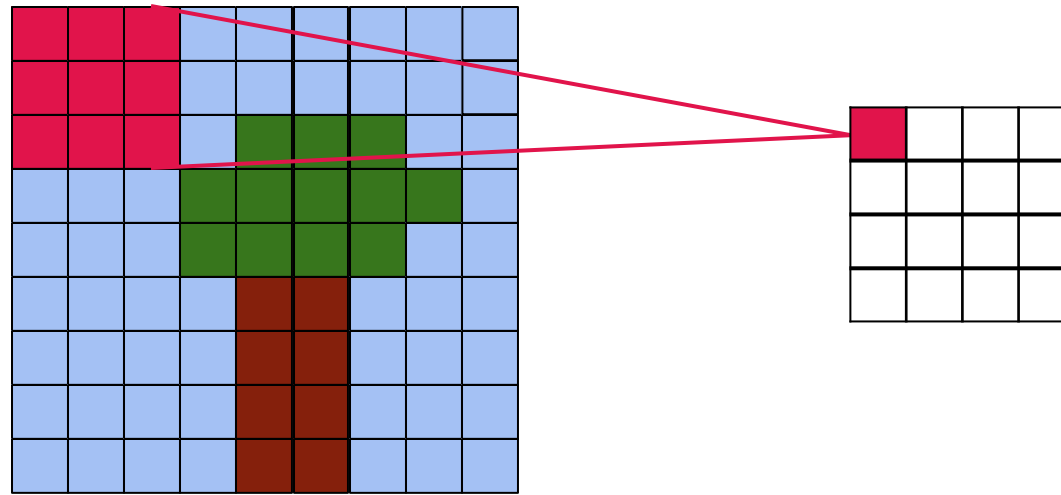
In the simplest case, the output value of the layer with input size  $(N, C_{\text{in}}, H, W)$  and output  $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$  can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$

- **out\_channels**
  - The number of different kernels;
  - Increase to detect different *types* of local patterns
  - Each out\_channel is a summation over input channels
- **kernel\_size**
  - The size of the kernel e.g., (3,3)
  - Increase to capture more *global* patterns

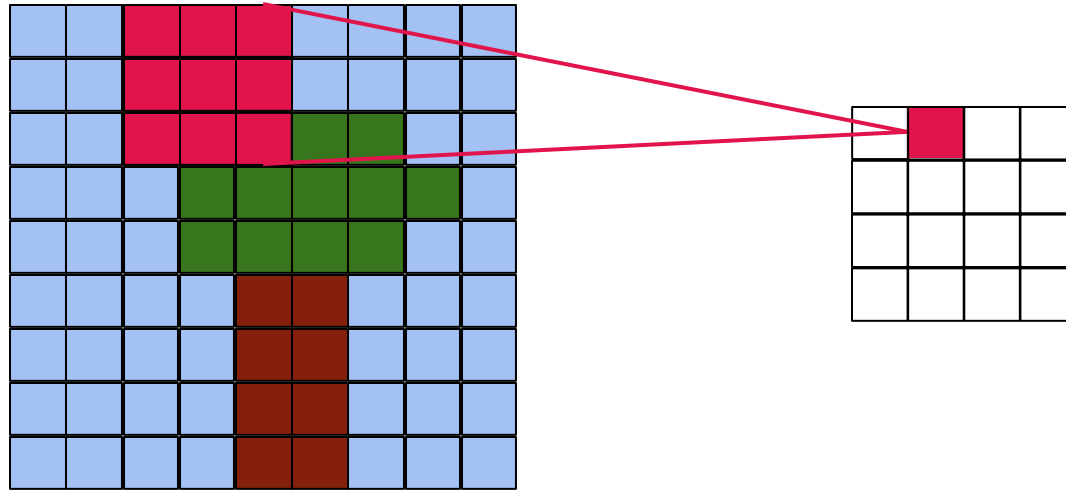
- **stride**
  - ?
- **padding**
  - ?
- **dilation**
  - ?

# Defining convolutional kernels: Stride



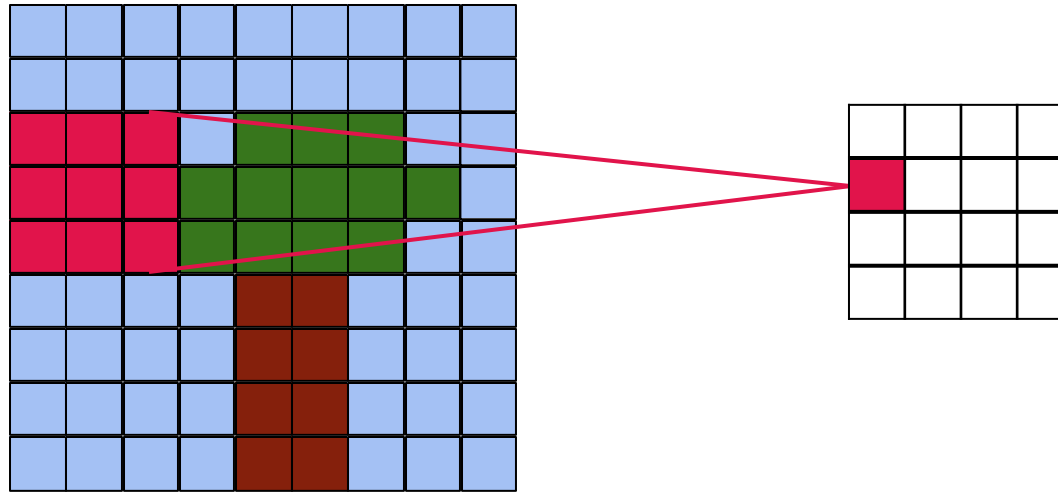
Stride = 2  $\Rightarrow$  kernel is applied every other pixel

# Defining convolutional kernels: Stride



Stride = 2  $\Rightarrow$  kernel is applied every other pixel

# Defining convolutional kernels: Stride



Stride = 2  $\Rightarrow$  kernel is applied every other pixel

# Defining convolutional kernels

## Conv2d

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1,
groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) [SOURCE]
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C_{in}, H, W)$  and output  $(N, C_{out}, H_{out}, W_{out})$  can be precisely described as:

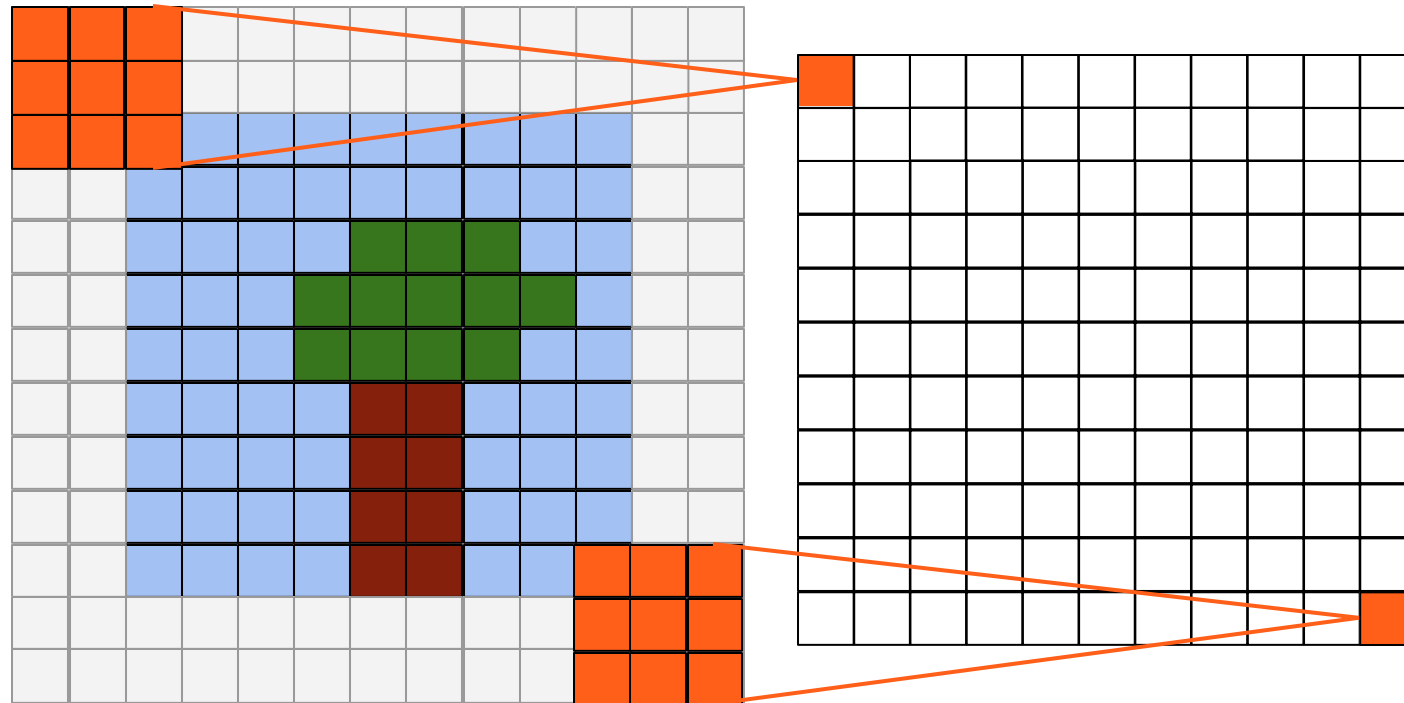
$$\text{out}(N_i, C_{out_j}) = \text{bias}(C_{out_j}) + \sum_{k=0}^{C_{in}-1} \text{weight}(C_{out_j}, k) \star \text{input}(N_i, k)$$

- **out\_channels**
  - The number of different kernels;
  - Increase to detect different *types* of local patterns
  - Each out\_channel is a summation over input channels
- **kernel\_size**
  - The size of the kernel e.g., (3,3)
  - Increase to capture more *global* patterns

- **stride**
  - Defines the adjacency of consecutive kernel applications
  - Enables the output to be further down sampled without increasing the kernel size
- **padding**
  - ?
- **dilation**
  - ?

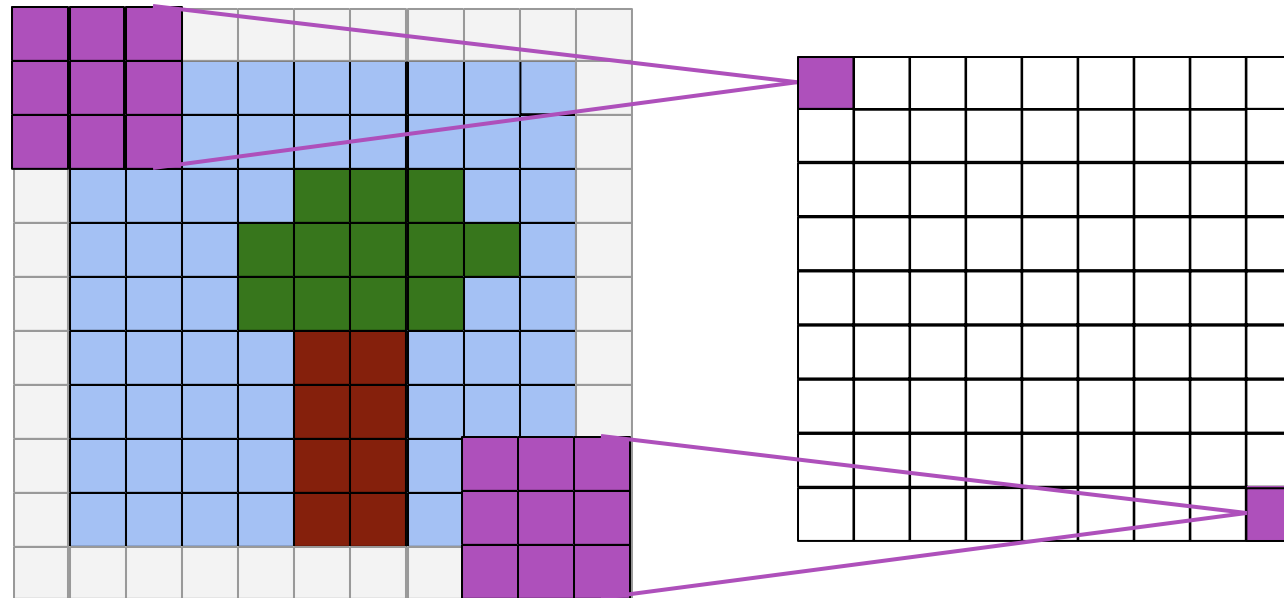


# Defining convolutional kernels: Padding



Padding = 2  $\Rightarrow$  dimensions are increased by 2 (of a provided value, generally 0). When combining padding of 2 and a kernel size of (3,3), the output dimension = input dimension + 1 (up sampling)

# Defining convolutional kernels: Padding



Padding = 1  $\Rightarrow$  dimensions are increased by 1 (of a provided value, generally 0). When combining padding of 1 and a kernel size of (3,3), the output dimension = input dimension

# Defining convolutional kernels

## Conv2d

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1,
groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) [SOURCE]
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C_{in}, H, W)$  and output  $(N, C_{out}, H_{out}, W_{out})$  can be precisely described as:

$$\text{out}(N_i, C_{out_j}) = \text{bias}(C_{out_j}) + \sum_{k=0}^{C_{in}-1} \text{weight}(C_{out_j}, k) \star \text{input}(N_i, k)$$

- **out\_channels**

- The number of different kernels;
- Increase to detect different *types* of local patterns
- Each out\_channel is a summation over input channels

- **kernel\_size**

- The size of the kernel e.g., (3,3)
- Increase to capture more *global* patterns

- **stride**

- Defines the adjacency of consecutive kernel applications
- Enables the output to be further down sampled without increasing the kernel size

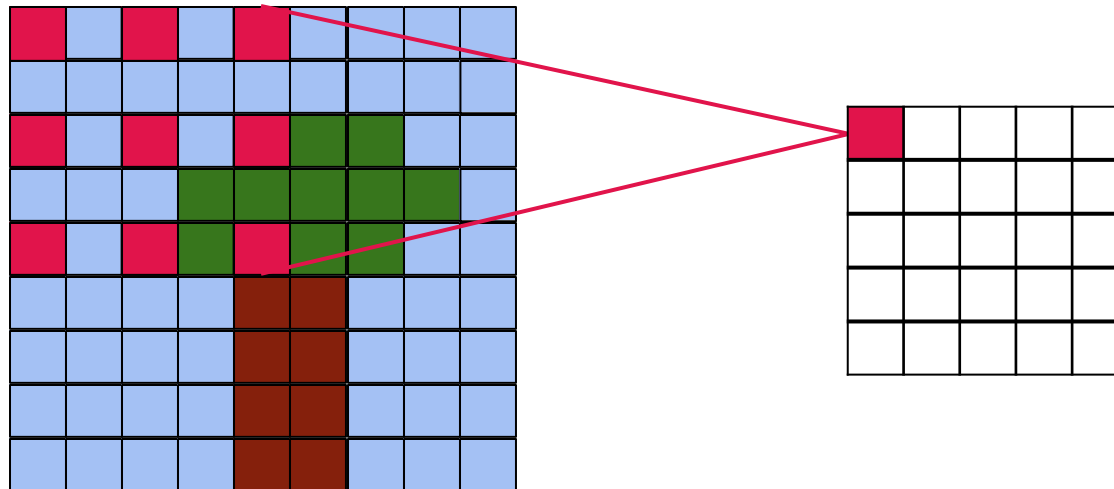
- **padding**

- Adds additional values to each dimension
- Manipulates the output dimensions (prevent down sampling\induce up sampling)

- **dilation**

- ?

# Defining convolutional kernels: Dilation



Dilation = 2  $\Rightarrow$  0 value, non-tunable weights of value 0 are interspersed every other weight. When combining dilation of 2 and a kernel size of (3,3), the kernel has *approximately* the same locality as a (5,5) kernel but with less parameters

# Defining convolutional kernels

## Conv2d

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1,
groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) \[SOURCE\]
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C_{in}, H, W)$  and output  $(N, C_{out}, H_{out}, W_{out})$  can be precisely described as:

$$\text{out}(N_i, C_{out_j}) = \text{bias}(C_{out_j}) + \sum_{k=0}^{C_{in}-1} \text{weight}(C_{out_j}, k) \star \text{input}(N_i, k)$$

- **out\_channels**

- The number of different kernels;
- Increase to detect different *types* of local patterns
- Each out\_channel is a summation over input channels

- **kernel\_size**

- The size of the kernel e.g., (3,3)
- Increase to capture more *global* patterns

- **stride**

- Defines the adjacency of consecutive kernel applications
- Enables the output to be further down sampled without increasing the kernel size

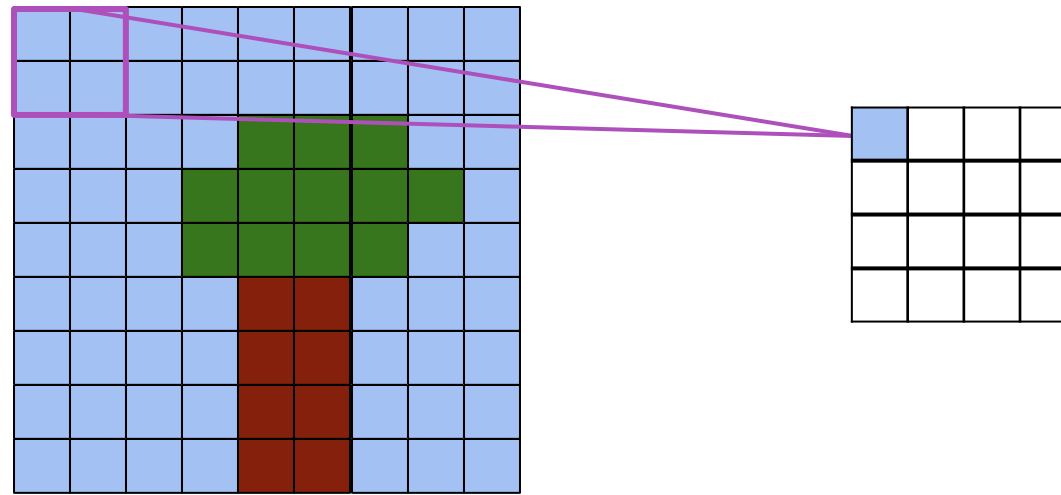
- **padding**

- Adds additional values to each dimension
- Manipulates the output dimensions (prevent down sampling\induce up sampling)

- **dilation**

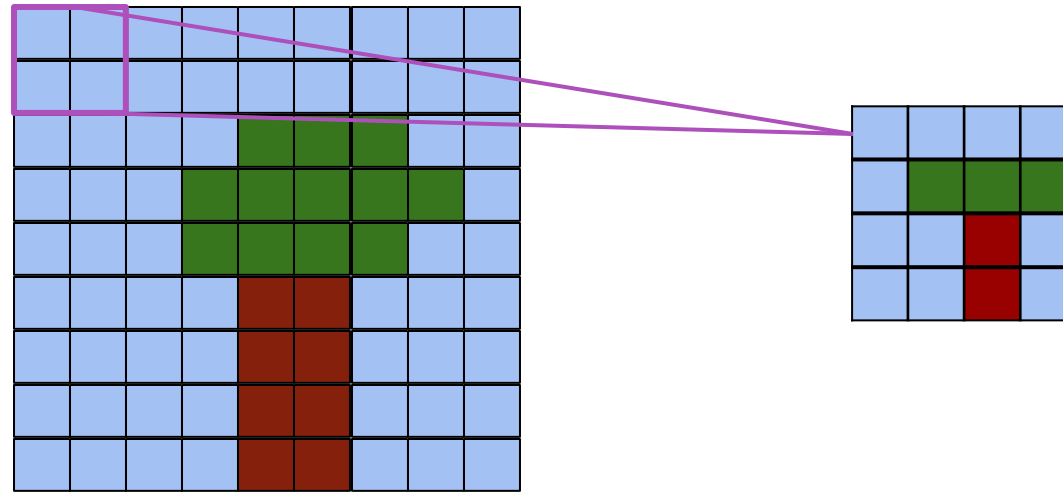
- Adds non-tunable weights to the filter
- Allows for larger locality without adding additional parameters

# Local pooling



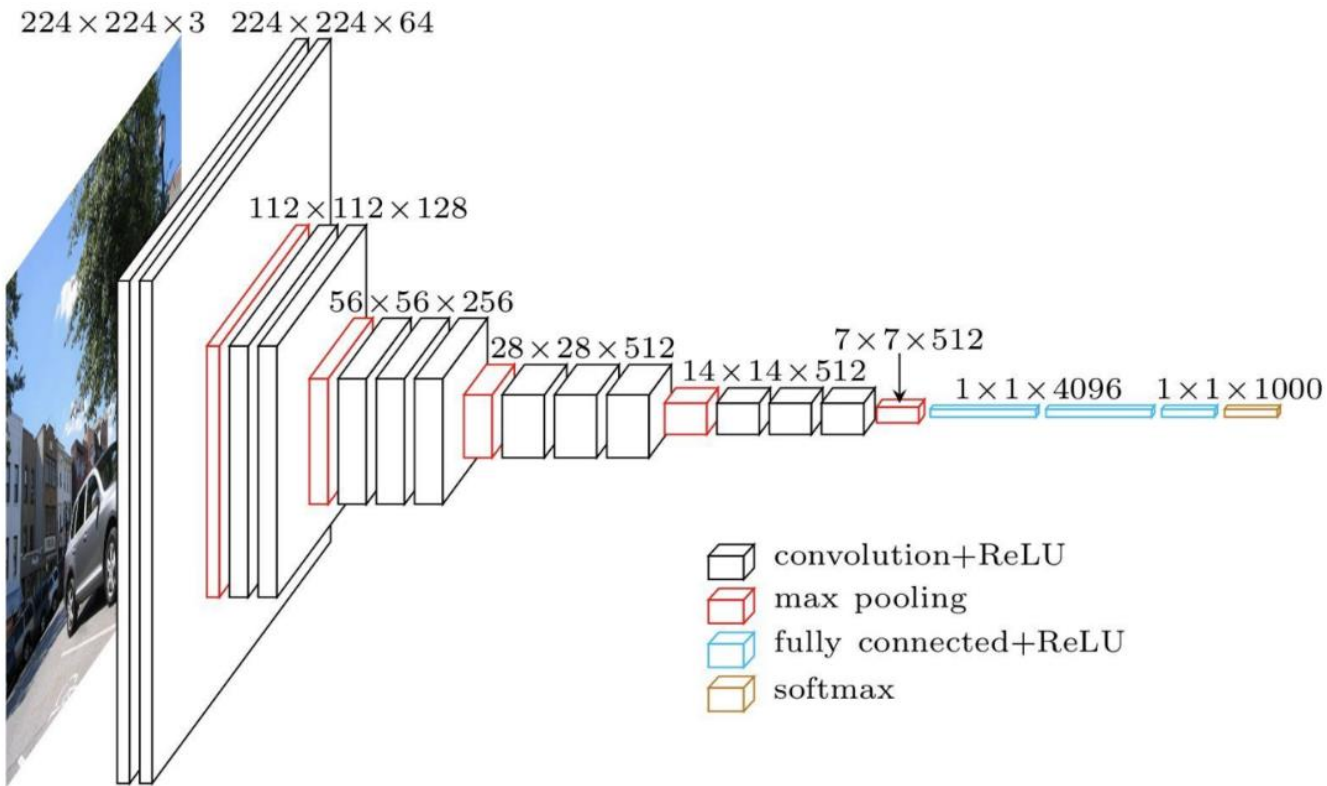
Computes summary statistics e.g, max and mean over filters (rather than convolving). Reduces resolution/coarsens the image

# Local pooling



Computes summary statistics e.g, max and mean over filters (rather than convolving). Reduces resolution/coarsens the image

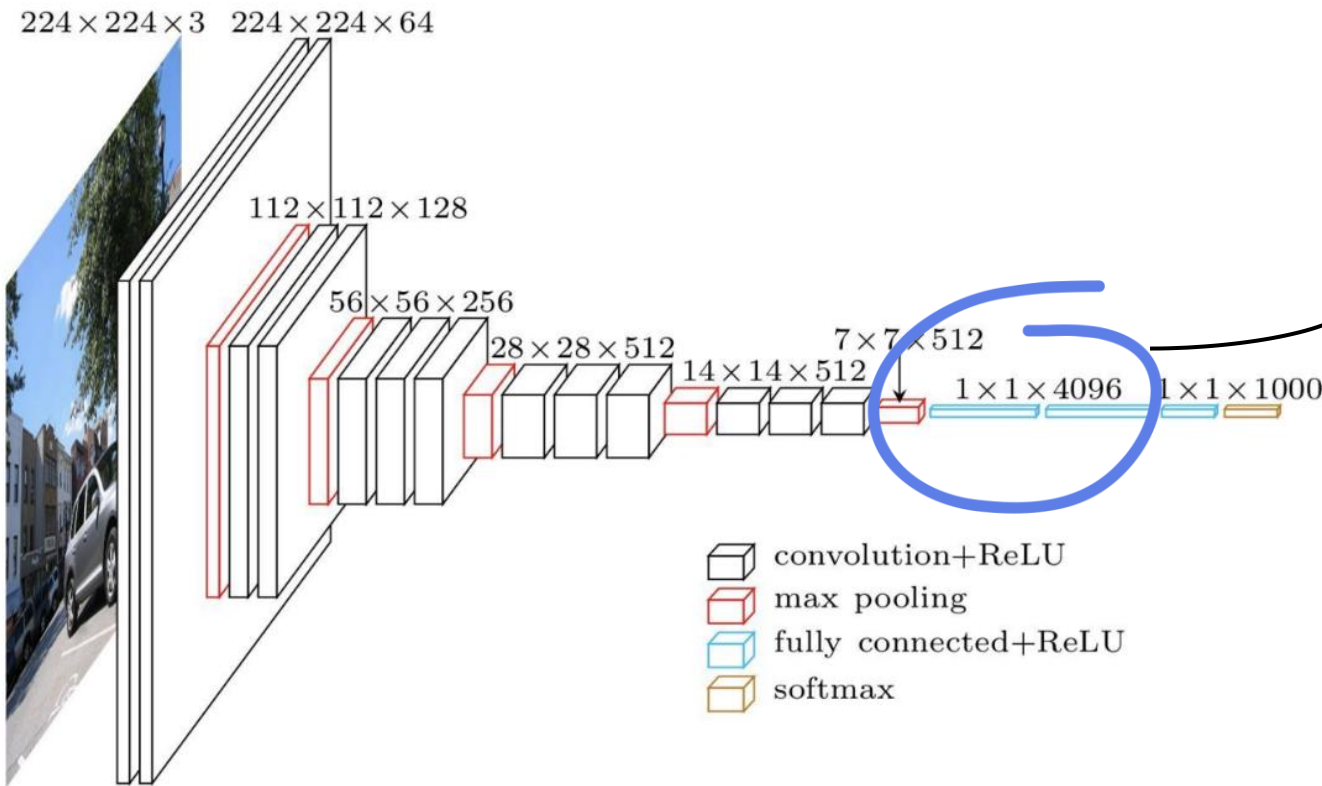
# VGG



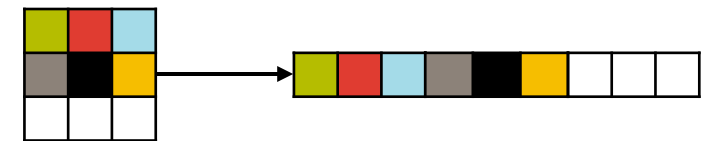
- Uses small kernels (3,3) with non-linear ReLU activations between conv layers
- Hugely upscales the number of filters whilst down sampling the image



# To MLPs (flatten/global pooling)



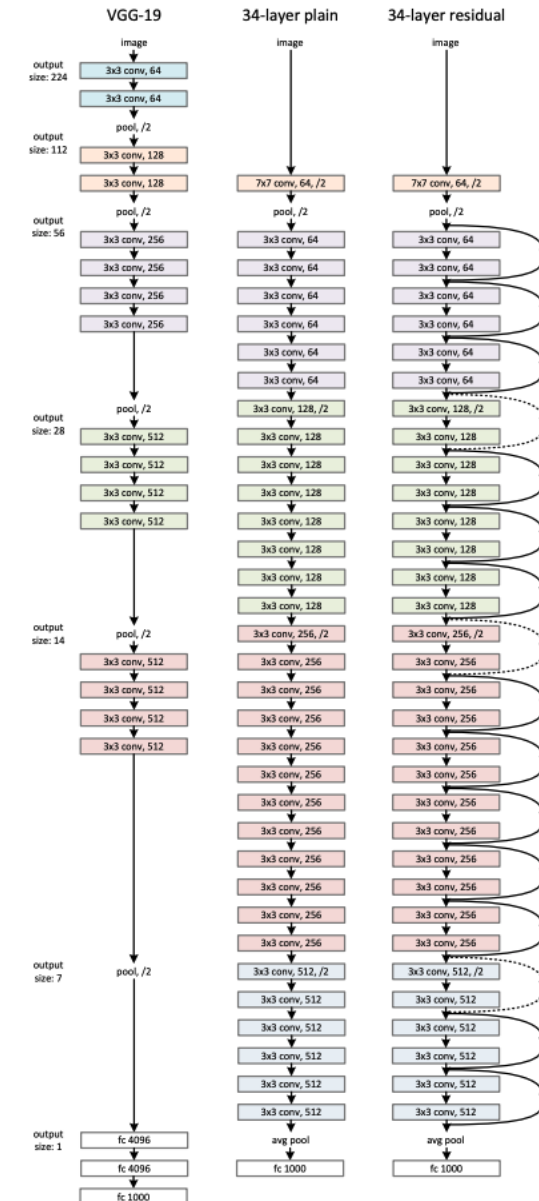
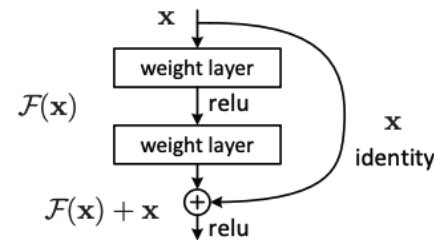
- MLPs require 1D input however, output of convolutions/pooling is 3D (7,7,512)
- Generally, “flattening” is used i.e.,



- Channel wise global pooling can also be used to reduce the dimensionality further however, this is not commonly used
  - Also induces invariance

# Deeper CNNs: ResNet

- Performance of deep networks (>10 layers) can saturate
  - Even when using batch norm/correct initialisation etc.
- ResNet proposes modelling
  - $g(z) = x$  where  $z = \alpha + \beta x$
  - The **residual** change is hypothesised to be easier to model than the full transformation



# Learning invariance and equivariance

# Learning invariance and equivariance

- Convolutions enable an **analytic expression of translational equivariance**
- When detecting objects **other desirable invariances and equivariances** included:
  - **Rotational**
    - Bounding boxes should equally rotate (equivariance)
    - Classification results should be the same (invariance)
  - **Scale**
    - Bounding boxes should become larger/smaller (equivariance)
    - Classification results should be the same (invariance)

# Data augmentations

- When there is no analytic expression, it might be possible to **augment** the input data with **simulated samples** to ensure that the model **learns the equ/invariance**
  - Images can be easily **rotated/cropped** and **scaled**
  - **Noise** can be overlayed to mimic 'poor quality' images or **parts of the image can be obscured**
- Not all domains are **as** amenable to data augmentations e.g., healthcare data
- Data augmentations are key for contrastive self-supervised learning of images (studied in week 5)

# Anonymous feedback

- <https://forms.gle/c3fFtxGG1aAUSrU56>